

The order of growth of the running time of an algorithm, defined in Chapter 2, gives a simple way to characterize the algorithm's efficiency and also allows us to compare it with alternative algorithms. Once the input size  $n$  becomes large enough, merge sort, with its  $\Theta(n \lg n)$  worst-case running time, beats insertion sort, whose worst-case running time is  $\Theta(n^2)$ . Although we can sometimes determine the exact running time of an algorithm, as we did for insertion sort in Chapter 2, the extra precision is rarely worth the effort of computing it. For large enough inputs, the multiplicative constants and lower-order terms of an exact running time are dominated by the effects of the input size itself.

When we look at input sizes large enough to make relevant only the order of growth of the running time, we are studying the *asymptotic* efficiency of algorithms. That is, we are concerned with how the running time of an algorithm increases with the size of the input *in the limit*, as the size of the input increases without bound. Usually, an algorithm that is asymptotically more efficient is the best choice for all but very small inputs.

This chapter gives several standard methods for simplifying the asymptotic analysis of algorithms. The next section presents informally the three most commonly used types of “asymptotic notation,” of which we have already seen an example in  $\Theta$ -notation. It also shows one way to use these asymptotic notations to reason about the worst-case running time of insertion sort. Then we look at asymptotic notations more formally and present several notational conventions used throughout this book. The last section reviews the behavior of functions that commonly arise when analyzing algorithms.

### 3.1 $O$ -notation, $\Omega$ -notation, and $\Theta$ -notation

When we analyzed the worst-case running time of insertion sort in Chapter 2, we started with the complicated expression

$$\left(\frac{c_5}{2} + \frac{c_6}{2} + \frac{c_7}{2}\right)n^2 + \left(c_1 + c_2 + c_4 + \frac{c_5}{2} - \frac{c_6}{2} - \frac{c_7}{2} + c_8\right)n - (c_2 + c_4 + c_5 + c_8).$$

We then discarded the lower-order terms  $(c_1 + c_2 + c_4 + c_5/2 - c_6/2 - c_7/2 + c_8)n$  and  $c_2 + c_4 + c_5 + c_8$ , and we also ignored the coefficient  $c_5/2 + c_6/2 + c_7/2$  of  $n^2$ . That left just the factor  $n^2$ , which we put into  $\Theta$ -notation as  $\Theta(n^2)$ . We use this style to characterize running times of algorithms: discard the lower-order terms and the coefficient of the leading term, and use a notation that focuses on the rate of growth of the running time.

$\Theta$ -notation is not the only such “asymptotic notation.” In this section, we’ll see other forms of asymptotic notation as well. We start with intuitive looks at these notations, revisiting insertion sort to see how we can apply them. In the next section, we’ll see the formal definitions of our asymptotic notations, along with conventions for using them.

Before we get into specifics, bear in mind that the asymptotic notations we’ll see are designed so that they characterize functions in general. It so happens that the functions we are most interested in denote the running times of algorithms. But asymptotic notation can apply to functions that characterize some other aspect of algorithms (the amount of space they use, for example), or even to functions that have nothing whatsoever to do with algorithms.

#### $O$ -notation

$O$ -notation characterizes an *upper bound* on the asymptotic behavior of a function. In other words, it says that a function grows *no faster* than a certain rate, based on the highest-order term. Consider, for example, the function  $7n^3 + 100n^2 - 20n + 6$ . Its highest-order term is  $7n^3$ , and so we say that this function’s rate of growth is  $n^3$ . Because this function grows no faster than  $n^3$ , we can write that it is  $O(n^3)$ . You might be surprised that we can also write that the function  $7n^3 + 100n^2 - 20n + 6$  is  $O(n^4)$ . Why? Because the function grows more slowly than  $n^4$ , we are correct in saying that it grows no faster. As you might have guessed, this function is also  $O(n^5)$ ,  $O(n^6)$ , and so on. More generally, it is  $O(n^c)$  for any constant  $c \geq 3$ .

**$\Omega$ -notation**

$\Omega$ -notation characterizes a *lower bound* on the asymptotic behavior of a function. In other words, it says that a function grows *at least as fast* as a certain rate, based—as in  $O$ -notation—on the highest-order term. Because the highest-order term in the function  $7n^3 + 100n^2 - 20n + 6$  grows at least as fast as  $n^3$ , this function is  $\Omega(n^3)$ . This function is also  $\Omega(n^2)$  and  $\Omega(n)$ . More generally, it is  $\Omega(n^c)$  for any constant  $c \leq 3$ .

 **$\Theta$ -notation**

$\Theta$ -notation characterizes a *tight bound* on the asymptotic behavior of a function. It says that a function grows *precisely* at a certain rate, based—once again—on the highest-order term. Put another way,  $\Theta$ -notation characterizes the rate of growth of the function to within a constant factor from above and to within a constant factor from below. These two constant factors need not be equal.

If you can show that a function is both  $O(f(n))$  and  $\Omega(f(n))$  for some function  $f(n)$ , then you have shown that the function is  $\Theta(f(n))$ . (The next section states this fact as a theorem.) For example, since the function  $7n^3 + 100n^2 - 20n + 6$  is both  $O(n^3)$  and  $\Omega(n^3)$ , it is also  $\Theta(n^3)$ .

**Example: Insertion sort**


Let's revisit insertion sort and see how to work with asymptotic notation to characterize its  $\Theta(n^2)$  worst-case running time without evaluating summations as we did in Chapter 2. Here is the INSERTION-SORT procedure once again:

```

INSERTION-SORT( $A, n$ )
1  for  $i = 2$  to  $n$ 
2       $key = A[i]$ 
3      // Insert  $A[i]$  into the sorted subarray  $A[1 : i - 1]$ .
4       $j = i - 1$ 
5      while  $j > 0$  and  $A[j] > key$ 
6           $A[j + 1] = A[j]$ 
7           $j = j - 1$ 
8       $A[j + 1] = key$ 

```

What can we observe about how the pseudocode operates? The procedure has nested loops. The outer loop is a **for** loop that runs  $n - 1$  times, regardless of the values being sorted. The inner loop is a **while** loop, but the number of iterations it makes depends on the values being sorted. The loop variable  $j$  starts at  $i - 1$



**Figure 3.1** The  $\Omega(n^2)$  lower bound for insertion sort. If the first  $n/3$  positions contain the  $n/3$  largest values, each of these values must move through each of the middle  $n/3$  positions, one position at a time, to end up somewhere in the last  $n/3$  positions. Since each of  $n/3$  values moves through at least each of  $n/3$  positions, the time taken in this case is at least proportional to  $(n/3)(n/3) = n^2/9$ , or  $\Omega(n^2)$ .

and decreases by 1 in each iteration until either it reaches 0 or  $A[j] \leq \text{key}$ . For a given value of  $i$ , the **while** loop might iterate 0 times,  $i - 1$  times, or anywhere in between. The body of the **while** loop (lines 6–7) takes constant time per iteration of the **while** loop.

These observations suffice to deduce an  $O(n^2)$  running time for any case of INSERTION-SORT, giving us a blanket statement that covers all inputs. The running time is dominated by the inner loop. Because each of the  $n - 1$  iterations of the outer loop causes the inner loop to iterate at most  $i - 1$  times, and because  $i$  is at most  $n$ , the total number of iterations of the inner loop is at most  $(n - 1)(n - 1)$ , which is less than  $n^2$ . Since each iteration of the inner loop takes constant time, the total time spent in the inner loop is at most a constant times  $n^2$ , or  $O(n^2)$ .

With a little creativity, we can also see that the worst-case running time of INSERTION-SORT is  $\Omega(n^2)$ . By saying that the worst-case running time of an algorithm is  $\Omega(n^2)$ , we mean that for every input size  $n$  above a certain threshold, there is at least one input of size  $n$  for which the algorithm takes at least  $cn^2$  time, for some positive constant  $c$ . It does not necessarily mean that the algorithm takes at least  $cn^2$  time for all inputs.

Let's now see why the worst-case running time of INSERTION-SORT is  $\Omega(n^2)$ . For a value to end up to the right of where it started, it must have been moved in line 6. In fact, for a value to end up  $k$  positions to the right of where it started, line 6 must have executed  $k$  times. As Figure 3.1 shows, let's assume that  $n$  is a multiple of 3 so that we can divide the array  $A$  into groups of  $n/3$  positions. Suppose that in the input to INSERTION-SORT, the  $n/3$  largest values occupy the first  $n/3$  array positions  $A[1 : n/3]$ . (It does not matter what relative order they have within the first  $n/3$  positions.) Once the array has been sorted, each of these  $n/3$  values ends up somewhere in the last  $n/3$  positions  $A[2n/3 + 1 : n]$ . For that to happen, each of these  $n/3$  values must pass through each of the middle  $n/3$  positions  $A[n/3 + 1 : 2n/3]$ . Each of these  $n/3$  values passes through these middle

$n/3$  positions one position at a time, by at least  $n/3$  executions of line 6. Because at least  $n/3$  values have to pass through at least  $n/3$  positions, the time taken by INSERTION-SORT in the worst case is at least proportional to  $(n/3)(n/3) = n^2/9$ , which is  $\Omega(n^2)$ .

Because we have shown that INSERTION-SORT runs in  $O(n^2)$  time in all cases and that there is an input that makes it take  $\Omega(n^2)$  time, we can conclude that the worst-case running time of INSERTION-SORT is  $\Theta(n^2)$ . It does not matter that the constant factors for upper and lower bounds might differ. What matters is that we have characterized the worst-case running time to within constant factors (discounting lower-order terms). This argument does not show that INSERTION-SORT runs in  $\Theta(n^2)$  time in *all* cases. Indeed, we saw in Chapter 2 that the best-case running time is  $\Theta(n)$ .

## Exercises

### 3.1-1

Modify the lower-bound argument for insertion sort to handle input sizes that are not necessarily a multiple of 3.

### 3.1-2

Using reasoning similar to what we used for insertion sort, analyze the running time of the selection sort algorithm from Exercise 2.2-2.


### 3.1-3

Suppose that  $\alpha$  is a fraction in the range  $0 < \alpha < 1$ . Show how to generalize the lower-bound argument for insertion sort to consider an input in which the  $\alpha n$  largest values start in the first  $\alpha n$  positions. What additional restriction do you need to put on  $\alpha$ ? What value of  $\alpha$  maximizes the number of times that the  $\alpha n$  largest values must pass through each of the middle  $(1 - 2\alpha)n$  array positions?

---

## 3.2 Asymptotic notation: formal definitions

Having seen asymptotic notation informally, let's get more formal. The notations we use to describe the asymptotic running time of an algorithm are defined in terms of functions whose domains are typically the set  $\mathbb{N}$  of natural numbers or the set  $\mathbb{R}$  of real numbers. Such notations are convenient for describing a running-time function  $T(n)$ . This section defines the basic asymptotic notations and also introduces some common “proper” notational abuses.



**Figure 3.2** Graphic examples of the  $O$ ,  $\Omega$ , and  $\Theta$  notations. In each part, the value of  $n_0$  shown is the minimum possible value, but any greater value also works. **(a)**  $O$ -notation gives an upper bound for a function to within a constant factor. We write  $f(n) = O(g(n))$  if there are positive constants  $n_0$  and  $c$  such that at and to the right of  $n_0$ , the value of  $f(n)$  always lies on or below  $cg(n)$ . **(b)**  $\Omega$ -notation gives a lower bound for a function to within a constant factor. We write  $f(n) = \Omega(g(n))$  if there are positive constants  $n_0$  and  $c$  such that at and to the right of  $n_0$ , the value of  $f(n)$  always lies on or above  $cg(n)$ . **(c)**  $\Theta$ -notation bounds a function to within constant factors. We write  $f(n) = \Theta(g(n))$  if there exist positive constants  $n_0$ ,  $c_1$ , and  $c_2$  such that at and to the right of  $n_0$ , the value of  $f(n)$  always lies between  $c_1g(n)$  and  $c_2g(n)$  inclusive.

### **$O$ -notation**

As we saw in Section 3.1,  $O$ -notation describes an *asymptotic upper bound*. We use  $O$ -notation to give an upper bound on a function, to within a constant factor.

Here is the formal definition of  $O$ -notation. For a given function  $g(n)$ , we denote by  $O(g(n))$  (pronounced “big-oh of  $g$  of  $n$ ” or sometimes just “oh of  $g$  of  $n$ ”) the *set of functions*

$$O(g(n)) = \{f(n) : \text{there exist positive constants } c \text{ and } n_0 \text{ such that } 0 \leq f(n) \leq cg(n) \text{ for all } n \geq n_0\}^1$$

A function  $f(n)$  belongs to the set  $O(g(n))$  if there exists a positive constant  $c$  such that  $f(n) \leq cg(n)$  for sufficiently large  $n$ . Figure 3.2(a) shows the intuition behind  $O$ -notation. For all values  $n$  at and to the right of  $n_0$ , the value of the function  $f(n)$  is on or below  $cg(n)$ .

The definition of  $O(g(n))$  requires that every function  $f(n)$  in the set  $O(g(n))$  be *asymptotically nonnegative*:  $f(n)$  must be nonnegative whenever  $n$  is sufficiently large. (An *asymptotically positive* function is one that is positive for all

<sup>1</sup> Within set notation, a colon means “such that.”

sufficiently large  $n$ .) Consequently, the function  $g(n)$  itself must be asymptotically nonnegative, or else the set  $O(g(n))$  is empty. We therefore assume that every function used within  $O$ -notation is asymptotically nonnegative. This assumption holds for the other asymptotic notations defined in this chapter as well.

You might be surprised that we define  $O$ -notation in terms of sets. Indeed, you might expect that we would write “ $f(n) \in O(g(n))$ ” to indicate that  $f(n)$  belongs to the set  $O(g(n))$ . Instead, we usually write “ $f(n) = O(g(n))$ ” and say “ $f(n)$  is big-oh of  $g(n)$ ” to express the same notion. Although it may seem confusing at first to abuse equality in this way, we’ll see later in this section that doing so has its advantages.

Let’s explore an example of how to use the formal definition of  $O$ -notation to justify our practice of discarding lower-order terms and ignoring the constant coefficient of the highest-order term. We’ll show that  $4n^2 + 100n + 500 = O(n^2)$ , even though the lower-order terms have much larger coefficients than the leading term. We need to find positive constants  $c$  and  $n_0$  such that  $4n^2 + 100n + 500 \leq cn^2$  for all  $n \geq n_0$ . Dividing both sides by  $n^2$  gives  $4 + 100/n + 500/n^2 \leq c$ . This inequality is satisfied for many choices of  $c$  and  $n_0$ . For example, if we choose  $n_0 = 1$ , then this inequality holds for  $c = 604$ . If we choose  $n_0 = 10$ , then  $c = 19$  works, and choosing  $n_0 = 100$  allows us to use  $c = 5.05$ .

We can also use the formal definition of  $O$ -notation to show that the function  $n^3 - 100n^2$  does not belong to the set  $O(n^2)$ , even though the coefficient of  $n^2$  is a large negative number. If we had  $n^3 - 100n^2 = O(n^2)$ , then there would be positive constants  $c$  and  $n_0$  such that  $n^3 - 100n^2 \leq cn^2$  for all  $n \geq n_0$ . Again, we divide both sides by  $n^2$ , giving  $n - 100 \leq c$ . Regardless of what value we choose for the constant  $c$ , this inequality does not hold for any value of  $n > c + 100$ .

### $\Omega$ -notation

Just as  $O$ -notation provides an asymptotic *upper* bound on a function,  $\Omega$ -notation provides an *asymptotic lower bound*. For a given function  $g(n)$ , we denote by  $\Omega(g(n))$  (pronounced “big-omega of  $g$  of  $n$ ” or sometimes just “omega of  $g$  of  $n$ ”) the set of functions

$$\Omega(g(n)) = \{f(n) : \text{there exist positive constants } c \text{ and } n_0 \text{ such that} \\ 0 \leq cg(n) \leq f(n) \text{ for all } n \geq n_0\}.$$

Figure 3.2(b) shows the intuition behind  $\Omega$ -notation. For all values  $n$  at or to the right of  $n_0$ , the value of  $f(n)$  is on or above  $cg(n)$ .

We’ve already shown that  $4n^2 + 100n + 500 = O(n^2)$ . Now let’s show that  $4n^2 + 100n + 500 = \Omega(n^2)$ . We need to find positive constants  $c$  and  $n_0$  such that  $4n^2 + 100n + 500 \geq cn^2$  for all  $n \geq n_0$ . As before, we divide both sides by  $n^2$ ,

giving  $4 + 100/n + 500/n^2 \geq c$ . This inequality holds when  $n_0$  is any positive integer and  $c = 4$ .

What if we had subtracted the lower-order terms from the  $4n^2$  term instead of adding them? What if we had a small coefficient for the  $n^2$  term? The function would still be  $\Omega(n^2)$ . For example, let's show that  $n^2/100 - 100n - 500 = \Omega(n^2)$ . Dividing by  $n^2$  gives  $1/100 - 100/n - 500/n^2 \geq c$ . We can choose any value for  $n_0$  that is at least 10,005 and find a positive value for  $c$ . For example, when  $n_0 = 10,005$ , we can choose  $c = 2.49 \times 10^{-9}$ . Yes, that's a tiny value for  $c$ , but it is positive. If we select a larger value for  $n_0$ , we can also increase  $c$ . For example, if  $n_0 = 100,000$ , then we can choose  $c = 0.0089$ . The higher the value of  $n_0$ , the closer to the coefficient  $1/100$  we can choose  $c$ .

### $\Theta$ -notation

We use  $\Theta$ -notation for *asymptotically tight bounds*. For a given function  $g(n)$ , we denote by  $\Theta(g(n))$  (“theta of  $g$  of  $n$ ”) the set of functions

$$\Theta(g(n)) = \{f(n) : \text{there exist positive constants } c_1, c_2, \text{ and } n_0 \text{ such that} \\ 0 \leq c_1 g(n) \leq f(n) \leq c_2 g(n) \text{ for all } n \geq n_0\}.$$

Figure 3.2(c) shows the intuition behind  $\Theta$ -notation. For all values of  $n$  at and to the right of  $n_0$ , the value of  $f(n)$  lies at or above  $c_1 g(n)$  and at or below  $c_2 g(n)$ . In other words, for all  $n \geq n_0$ , the function  $f(n)$  is equal to  $g(n)$  to within constant factors.

The definitions of  $O$ -,  $\Omega$ -, and  $\Theta$ -notations lead to the following theorem, whose proof we leave as Exercise 3.2-4.

#### **Theorem 3.1**

For any two functions  $f(n)$  and  $g(n)$ , we have  $f(n) = \Theta(g(n))$  if and only if  $f(n) = O(g(n))$  and  $f(n) = \Omega(g(n))$ . ■

We typically apply Theorem 3.1 to prove asymptotically tight bounds from asymptotic upper and lower bounds.

### Asymptotic notation and running times

When you use asymptotic notation to characterize an algorithm's running time, make sure that the asymptotic notation you use is as precise as possible without overstating which running time it applies to. Here are some examples of using asymptotic notation properly and improperly to characterize running times.

Let's start with insertion sort. We can correctly say that insertion sort's worst-case running time is  $O(n^2)$ ,  $\Omega(n^2)$ , and—due to Theorem 3.1— $\Theta(n^2)$ . Although



all three ways to characterize the worst-case running times are correct, the  $\Theta(n^2)$  bound is the most precise and hence the most preferred. We can also correctly say that insertion sort's best-case running time is  $O(n)$ ,  $\Omega(n)$ , and  $\Theta(n)$ , again with  $\Theta(n)$  the most precise and therefore the most preferred.

Here is what we *cannot* correctly say: insertion sort's running time is  $\Theta(n^2)$ . That is an overstatement because by omitting “worst-case” from the statement, we're left with a blanket statement covering all cases. The error here is that insertion sort does not run in  $\Theta(n^2)$  time in all cases since, as we've seen, it runs in  $\Theta(n)$  time in the best case. We can correctly say that insertion sort's running time is  $O(n^2)$ , however, because in all cases, its running time grows no faster than  $n^2$ . When we say  $O(n^2)$  instead of  $\Theta(n^2)$ , there is no problem in having cases whose running time grows more slowly than  $n^2$ . Likewise, we cannot correctly say that insertion sort's running time is  $\Theta(n)$ , but we can say that its running time is  $\Omega(n)$ .

How about merge sort? Since merge sort runs in  $\Theta(n \lg n)$  time in all cases, we can just say that its running time is  $\Theta(n \lg n)$  without specifying worst-case, best-case, or any other case.

People occasionally conflate  $O$ -notation with  $\Theta$ -notation by mistakenly using  $O$ -notation to indicate an asymptotically tight bound. They say things like “an  $O(n \lg n)$ -time algorithm runs faster than an  $O(n^2)$ -time algorithm.” Maybe it does, maybe it doesn't. Since  $O$ -notation denotes only an asymptotic upper bound, that so-called  $O(n^2)$ -time algorithm might actually run in  $\Theta(n)$  time. You should be careful to choose the appropriate asymptotic notation. If you want to indicate an asymptotically tight bound, use  $\Theta$ -notation.

We typically use asymptotic notation to provide the simplest and most precise bounds possible. For example, if an algorithm has a running time of  $3n^2 + 20n$  in all cases, we use asymptotic notation to write that its running time is  $\Theta(n^2)$ . Strictly speaking, we are also correct in writing that the running time is  $O(n^3)$  or  $\Theta(3n^2 + 20n)$ . Neither of these expressions is as useful as writing  $\Theta(n^2)$  in this case, however:  $O(n^3)$  is less precise than  $\Theta(n^2)$  if the running time is  $3n^2 + 20n$ , and  $\Theta(3n^2 + 20n)$  introduces complexity that obscures the order of growth. By writing the simplest and most precise bound, such as  $\Theta(n^2)$ , we can categorize and compare different algorithms. Throughout the book, you will see asymptotic running times that are almost always based on polynomials and logarithms: functions such as  $n$ ,  $n \lg^2 n$ ,  $n^2 \lg n$ , or  $n^{1/2}$ . You will also see some other functions, such as exponentials,  $\lg \lg n$ , and  $\lg^* n$  (see Section 3.3). It is usually fairly easy to compare the rates of growth of these functions. Problem 3-3 gives you good practice.

### Asymptotic notation in equations and inequalities

Although we formally define asymptotic notation in terms of sets, we use the equal sign ( $=$ ) instead of the set membership sign ( $\in$ ) within formulas. For example, we wrote that  $4n^2 + 100n + 500 = O(n^2)$ . We might also write  $2n^2 + 3n + 1 = 2n^2 + \Theta(n)$ . How do we interpret such formulas?

When the asymptotic notation stands alone (that is, not within a larger formula) on the right-hand side of an equation (or inequality), as in  $4n^2 + 100n + 500 = O(n^2)$ , the equal sign means set membership:  $4n^2 + 100n + 500 \in O(n^2)$ . In general, however, when asymptotic notation appears in a formula, we interpret it as standing for some anonymous function that we do not care to name. For example, the formula  $2n^2 + 3n + 1 = 2n^2 + \Theta(n)$  means that  $2n^2 + 3n + 1 = 2n^2 + f(n)$ , where  $f(n) \in \Theta(n)$ . In this case, we let  $f(n) = 3n + 1$ , which indeed belongs to  $\Theta(n)$ .

Using asymptotic notation in this manner can help eliminate inessential detail and clutter in an equation. For example, in Chapter 2 we expressed the worst-case running time of merge sort as the recurrence

$$T(n) = 2T(n/2) + \Theta(n) .$$

If we are interested only in the asymptotic behavior of  $T(n)$ , there is no point in specifying all the lower-order terms exactly, because they are all understood to be included in the anonymous function denoted by the term  $\Theta(n)$ .

The number of anonymous functions in an expression is understood to be equal to the number of times the asymptotic notation appears. For example, in the expression

$$\sum_{i=1}^n O(i) ,$$

there is only a single anonymous function (a function of  $i$ ). This expression is thus *not* the same as  $O(1) + O(2) + \cdots + O(n)$ , which doesn't really have a clean interpretation.

In some cases, asymptotic notation appears on the left-hand side of an equation, as in

$$2n^2 + \Theta(n) = \Theta(n^2) .$$

Interpret such equations using the following rule: *No matter how the anonymous functions are chosen on the left of the equal sign, there is a way to choose the anonymous functions on the right of the equal sign to make the equation valid.* Thus, our example means that for *any* function  $f(n) \in \Theta(n)$ , there is *some* function  $g(n) \in \Theta(n^2)$  such that  $2n^2 + f(n) = g(n)$  for all  $n$ . In other words, the right-hand side of an equation provides a coarser level of detail than the left-hand side.

We can chain together a number of such relationships, as in

$$\begin{aligned} 2n^2 + 3n + 1 &= 2n^2 + \Theta(n) \\ &= \Theta(n^2) . \end{aligned}$$

By the rules above, interpret each equation separately. The first equation says that there is *some* function  $f(n) \in \Theta(n)$  such that  $2n^2 + 3n + 1 = 2n^2 + f(n)$  for all  $n$ . The second equation says that for *any* function  $g(n) \in \Theta(n)$  (such as the  $f(n)$  just mentioned), there is *some* function  $h(n) \in \Theta(n^2)$  such that  $2n^2 + g(n) = h(n)$  for all  $n$ . This interpretation implies that  $2n^2 + 3n + 1 = \Theta(n^2)$ , which is what the chaining of equations intuitively says.

### Proper abuses of asymptotic notation

Besides the abuse of equality to mean set membership, which we now see has a precise mathematical interpretation, another abuse of asymptotic notation occurs when the variable tending toward  $\infty$  must be inferred from context. For example, when we say  $O(g(n))$ , we can assume that we're interested in the growth of  $g(n)$  as  $n$  grows, and if we say  $O(g(m))$  we're talking about the growth of  $g(m)$  as  $m$  grows. The free variable in the expression indicates what variable is going to  $\infty$ .

The most common situation requiring contextual knowledge of which variable tends to  $\infty$  occurs when the function inside the asymptotic notation is a constant, as in the expression  $O(1)$ . We cannot infer from the expression which variable is going to  $\infty$ , because no variable appears there. The context must disambiguate. For example, if the equation using asymptotic notation is  $f(n) = O(1)$ , it's apparent that the variable we're interested in is  $n$ . Knowing from context that the variable of interest is  $n$ , however, allows us to make perfect sense of the expression by using the formal definition of  $O$ -notation: the expression  $f(n) = O(1)$  means that the function  $f(n)$  is bounded from above by a constant as  $n$  goes to  $\infty$ . Technically, it might be less ambiguous if we explicitly indicated the variable tending to  $\infty$  in the asymptotic notation itself, but that would clutter the notation. Instead, we simply ensure that the context makes it clear which variable (or variables) tend to  $\infty$ .

When the function inside the asymptotic notation is bounded by a positive constant, as in  $T(n) = O(1)$ , we often abuse asymptotic notation in yet another way, especially when stating recurrences. We may write something like  $T(n) = O(1)$  for  $n < 3$ . According to the formal definition of  $O$ -notation, this statement is meaningless, because the definition only says that  $T(n)$  is bounded above by a positive constant  $c$  for  $n \geq n_0$  for some  $n_0 > 0$ . The value of  $T(n)$  for  $n < n_0$  need not be so bounded. Thus, in the example  $T(n) = O(1)$  for  $n < 3$ , we cannot infer any constraint on  $T(n)$  when  $n < 3$ , because it might be that  $n_0 > 3$ .

What is conventionally meant when we say  $T(n) = O(1)$  for  $n < 3$  is that there exists a positive constant  $c$  such that  $T(n) \leq c$  for  $n < 3$ . This convention saves

us the trouble of naming the bounding constant, allowing it to remain anonymous while we focus on more important variables in an analysis. Similar abuses occur with the other asymptotic notations. For example,  $T(n) = \Theta(1)$  for  $n < 3$  means that  $T(n)$  is bounded above and below by positive constants when  $n < 3$ .

Occasionally, the function describing an algorithm's running time may not be defined for certain input sizes, for example, when an algorithm assumes that the input size is an exact power of 2. We still use asymptotic notation to describe the growth of the running time, understanding that any constraints apply only when the function is defined. For example, suppose that  $f(n)$  is defined only on a subset of the natural or nonnegative real numbers. Then  $f(n) = O(g(n))$  means that the bound  $0 \leq T(n) \leq cg(n)$  in the definition of  $O$ -notation holds for all  $n \geq n_0$  over the domain of  $f(n)$ , that is, where  $f(n)$  is defined. This abuse is rarely pointed out, since what is meant is generally clear from context.

In mathematics, it's okay—and often desirable—to abuse a notation, as long as we don't misuse it. If we understand precisely what is meant by the abuse and don't draw incorrect conclusions, it can simplify our mathematical language, contribute to our higher-level understanding, and help us focus on what really matters.

### ***o*-notation**

The asymptotic upper bound provided by  $O$ -notation may or may not be asymptotically tight. The bound  $2n^2 = O(n^2)$  is asymptotically tight, but the bound  $2n = O(n^2)$  is not. We use  $o$ -notation to denote an upper bound that is not asymptotically tight. We formally define  $o(g(n))$  ("little-oh of  $g$  of  $n$ ") as the set

$$o(g(n)) = \{f(n) : \text{for any positive constant } c > 0, \text{ there exists a constant } n_0 > 0 \text{ such that } 0 \leq f(n) < cg(n) \text{ for all } n \geq n_0\}.$$

For example,  $2n = o(n^2)$ , but  $2n^2 \neq o(n^2)$ .

The definitions of  $O$ -notation and  $o$ -notation are similar. The main difference is that in  $f(n) = O(g(n))$ , the bound  $0 \leq f(n) \leq cg(n)$  holds for *some* constant  $c > 0$ , but in  $f(n) = o(g(n))$ , the bound  $0 \leq f(n) < cg(n)$  holds for *all* constants  $c > 0$ . Intuitively, in  $o$ -notation, the function  $f(n)$  becomes insignificant relative to  $g(n)$  as  $n$  gets large:

$$\lim_{n \rightarrow \infty} \frac{f(n)}{g(n)} = 0.$$

Some authors use this limit as a definition of the  $o$ -notation, but the definition in this book also restricts the anonymous functions to be asymptotically nonnegative.

**$\omega$ -notation**

By analogy,  $\omega$ -notation is to  $\Omega$ -notation as  $o$ -notation is to  $O$ -notation. We use  $\omega$ -notation to denote a lower bound that is not asymptotically tight. One way to define it is by

$f(n) \in \omega(g(n))$  if and only if  $g(n) \in o(f(n))$ .

Formally, however, we define  $\omega(g(n))$  (“little-omega of  $g$  of  $n$ ”) as the set

$$\omega(g(n)) = \{f(n) : \text{for any positive constant } c > 0, \text{ there exists a constant } n_0 > 0 \text{ such that } 0 \leq cg(n) < f(n) \text{ for all } n \geq n_0\}.$$

Where the definition of  $o$ -notation says that  $f(n) < cg(n)$ , the definition of  $\omega$ -notation says the opposite: that  $cg(n) < f(n)$ . For examples of  $\omega$ -notation, we have  $n^2/2 = \omega(n)$ , but  $n^2/2 \neq \omega(n^2)$ . The relation  $f(n) = \omega(g(n))$  implies that

$$\lim_{n \rightarrow \infty} \frac{f(n)}{g(n)} = \infty,$$

if the limit exists. That is,  $f(n)$  becomes arbitrarily large relative to  $g(n)$  as  $n$  gets large.

**Comparing functions**

Many of the relational properties of real numbers apply to asymptotic comparisons as well. For the following, assume that  $f(n)$  and  $g(n)$  are asymptotically positive.

**Transitivity:**

$$\begin{aligned} f(n) = \Theta(g(n)) \text{ and } g(n) = \Theta(h(n)) & \text{ imply } f(n) = \Theta(h(n)), \\ f(n) = O(g(n)) \text{ and } g(n) = O(h(n)) & \text{ imply } f(n) = O(h(n)), \\ f(n) = \Omega(g(n)) \text{ and } g(n) = \Omega(h(n)) & \text{ imply } f(n) = \Omega(h(n)), \\ f(n) = o(g(n)) \text{ and } g(n) = o(h(n)) & \text{ imply } f(n) = o(h(n)), \\ f(n) = \omega(g(n)) \text{ and } g(n) = \omega(h(n)) & \text{ imply } f(n) = \omega(h(n)). \end{aligned}$$

**Reflexivity:**

$$\begin{aligned} f(n) &= \Theta(f(n)), \\ f(n) &= O(f(n)), \\ f(n) &= \Omega(f(n)). \end{aligned}$$

**Symmetry:**

$$f(n) = \Theta(g(n)) \text{ if and only if } g(n) = \Theta(f(n)).$$

**Transpose symmetry:**

$$f(n) = O(g(n)) \text{ if and only if } g(n) = \Omega(f(n)),$$

$$f(n) = o(g(n)) \text{ if and only if } g(n) = \omega(f(n)).$$

Because these properties hold for asymptotic notations, we can draw an analogy between the asymptotic comparison of two functions  $f$  and  $g$  and the comparison of two real numbers  $a$  and  $b$ :

$$f(n) = O(g(n)) \text{ is like } a \leq b,$$

$$f(n) = \Omega(g(n)) \text{ is like } a \geq b,$$

$$f(n) = \Theta(g(n)) \text{ is like } a = b,$$

$$f(n) = o(g(n)) \text{ is like } a < b,$$

$$f(n) = \omega(g(n)) \text{ is like } a > b.$$

We say that  $f(n)$  is *asymptotically smaller* than  $g(n)$  if  $f(n) = o(g(n))$ , and  $f(n)$  is *asymptotically larger* than  $g(n)$  if  $f(n) = \omega(g(n))$ .

One property of real numbers, however, does not carry over to asymptotic notation:

**Trichotomy:** For any two real numbers  $a$  and  $b$ , exactly one of the following must hold:  $a < b$ ,  $a = b$ , or  $a > b$ .

Although any two real numbers can be compared, not all functions are asymptotically comparable. That is, for two functions  $f(n)$  and  $g(n)$ , it may be the case that neither  $f(n) = O(g(n))$  nor  $f(n) = \Omega(g(n))$  holds. For example, we cannot compare the functions  $n$  and  $n^{1+\sin n}$  using asymptotic notation, since the value of the exponent in  $n^{1+\sin n}$  oscillates between 0 and 2, taking on all values in between.

**Exercises****3.2-1**

Let  $f(n)$  and  $g(n)$  be asymptotically nonnegative functions. Using the basic definition of  $\Theta$ -notation, prove that  $\max\{f(n), g(n)\} = \Theta(f(n) + g(n))$ .

**3.2-2**

Explain why the statement, “The running time of algorithm  $A$  is at least  $O(n^2)$ ,” is meaningless.

**3.2-3**

Is  $2^{n+1} = O(2^n)$ ? Is  $2^{2n} = O(2^n)$ ?

**3.2-4**

Prove Theorem 3.1.

**3.2-5**

Prove that the running time of an algorithm is  $\Theta(g(n))$  if and only if its worst-case running time is  $O(g(n))$  and its best-case running time is  $\Omega(g(n))$ .

**3.2-6**

Prove that  $o(g(n)) \cap \omega(g(n))$  is the empty set.

**3.2-7**

We can extend our notation to the case of two parameters  $n$  and  $m$  that can go to  $\infty$  independently at different rates. For a given function  $g(n, m)$ , we denote by  $O(g(n, m))$  the set of functions

$$O(g(n, m)) = \{f(n, m) : \text{there exist positive constants } c, n_0, \text{ and } m_0 \\ \text{such that } 0 \leq f(n, m) \leq cg(n, m) \\ \text{for all } n \geq n_0 \text{ or } m \geq m_0\}.$$

Give corresponding definitions for  $\Omega(g(n, m))$  and  $\Theta(g(n, m))$ .

### 3.3 Standard notations and common functions

This section reviews some standard mathematical functions and notations and explores the relationships among them. It also illustrates the use of the asymptotic notations.

#### Monotonicity

A function  $f(n)$  is **monotonically increasing** if  $m \leq n$  implies  $f(m) \leq f(n)$ . Similarly, it is **monotonically decreasing** if  $m \leq n$  implies  $f(m) \geq f(n)$ . A function  $f(n)$  is **strictly increasing** if  $m < n$  implies  $f(m) < f(n)$  and **strictly decreasing** if  $m < n$  implies  $f(m) > f(n)$ .

#### Floors and ceilings

For any real number  $x$ , we denote the greatest integer less than or equal to  $x$  by  $\lfloor x \rfloor$  (read “the floor of  $x$ ”) and the least integer greater than or equal to  $x$  by  $\lceil x \rceil$  (read “the ceiling of  $x$ ”). The floor function is monotonically increasing, as is the ceiling function.

Floors and ceilings obey the following properties. For any integer  $n$ , we have

$$\lfloor n \rfloor = n = \lceil n \rceil. \quad (3.1)$$

For all real  $x$ , we have

$$x - 1 < \lfloor x \rfloor \leq x \leq \lceil x \rceil < x + 1 . \quad (3.2)$$

We also have

$$-\lfloor x \rfloor = \lceil -x \rceil , \quad (3.3)$$

or equivalently,

$$-\lceil x \rceil = \lfloor -x \rfloor . \quad (3.4)$$

For any real number  $x \geq 0$  and integers  $a, b > 0$ , we have

$$\left\lceil \frac{\lfloor x/a \rfloor}{b} \right\rceil = \left\lceil \frac{x}{ab} \right\rceil , \quad (3.5)$$

$$\left\lfloor \frac{\lceil x/a \rceil}{b} \right\rfloor = \left\lfloor \frac{x}{ab} \right\rfloor , \quad (3.6)$$

$$\left\lceil \frac{a}{b} \right\rceil \leq \frac{a + (b - 1)}{b} , \quad (3.7)$$

$$\left\lfloor \frac{a}{b} \right\rfloor \geq \frac{a - (b - 1)}{b} . \quad (3.8)$$

For any integer  $n$  and real number  $x$ , we have

$$\lfloor n + x \rfloor = n + \lfloor x \rfloor , \quad (3.9)$$

$$\lceil n + x \rceil = n + \lceil x \rceil . \quad (3.10)$$

### Modular arithmetic

For any integer  $a$  and any positive integer  $n$ , the value  $a \bmod n$  is the *remainder* (or *residue*) of the quotient  $a/n$ :

$$a \bmod n = a - n \lfloor a/n \rfloor . \quad (3.11)$$

It follows that

$$0 \leq a \bmod n < n , \quad (3.12)$$

even when  $a$  is negative.

Given a well-defined notion of the remainder of one integer when divided by another, it is convenient to provide special notation to indicate equality of remainders. If  $(a \bmod n) = (b \bmod n)$ , we write  $a = b \pmod{n}$  and say that  $a$  is *equivalent* to  $b$ , modulo  $n$ . In other words,  $a = b \pmod{n}$  if  $a$  and  $b$  have the same remainder when divided by  $n$ . Equivalently,  $a = b \pmod{n}$  if and only if  $n$  is a divisor of  $b - a$ . We write  $a \neq b \pmod{n}$  if  $a$  is not equivalent to  $b$ , modulo  $n$ .



## Polynomials

Given a nonnegative integer  $d$ , a **polynomial in  $n$  of degree  $d$**  is a function  $p(n)$  of the form

$$p(n) = \sum_{i=0}^d a_i n^i ,$$

where the constants  $a_0, a_1, \dots, a_d$  are the **coefficients** of the polynomial and  $a_d \neq 0$ . A polynomial is asymptotically positive if and only if  $a_d > 0$ . For an asymptotically positive polynomial  $p(n)$  of degree  $d$ , we have  $p(n) = \Theta(n^d)$ . For any real constant  $a \geq 0$ , the function  $n^a$  is monotonically increasing, and for any real constant  $a \leq 0$ , the function  $n^a$  is monotonically decreasing. We say that a function  $f(n)$  is **polynomially bounded** if  $f(n) = O(n^k)$  for some constant  $k$ .

## Exponentials

For all real  $a > 0$ ,  $m$ , and  $n$ , we have the following identities:

$$\begin{aligned} a^0 &= 1 , \\ a^1 &= a , \\ a^{-1} &= 1/a , \\ (a^m)^n &= a^{mn} , \\ (a^m)^n &= (a^n)^m , \\ a^m a^n &= a^{m+n} . \end{aligned}$$

For all  $n$  and  $a \geq 1$ , the function  $a^n$  is monotonically increasing in  $n$ . When convenient, we assume that  $0^0 = 1$ .

We can relate the rates of growth of polynomials and exponentials by the following fact. For all real constants  $a > 1$  and  $b$ , we have

$$\lim_{n \rightarrow \infty} \frac{n^b}{a^n} = 0 ,$$

from which we can conclude that

$$n^b = o(a^n) . \tag{3.13}$$

Thus, any exponential function with a base strictly greater than 1 grows faster than any polynomial function.

Using  $e$  to denote  $2.71828\dots$ , the base of the natural-logarithm function, we have for all real  $x$ ,

$$e^x = 1 + x + \frac{x^2}{2!} + \frac{x^3}{3!} + \dots = \sum_{i=0}^{\infty} \frac{x^i}{i!} ,$$

where “!” denotes the factorial function defined later in this section. For all real  $x$ , we have the inequality

$$1 + x \leq e^x, \quad (3.14)$$

where equality holds only when  $x = 0$ . When  $|x| \leq 1$ , we have the approximation

$$1 + x \leq e^x \leq 1 + x + x^2. \quad (3.15)$$

When  $x \rightarrow 0$ , the approximation of  $e^x$  by  $1 + x$  is quite good:

$$e^x = 1 + x + \Theta(x^2).$$

(In this equation, the asymptotic notation is used to describe the limiting behavior as  $x \rightarrow 0$  rather than as  $x \rightarrow \infty$ .) We have for all  $x$ ,

$$\lim_{n \rightarrow \infty} \left(1 + \frac{x}{n}\right)^n = e^x. \quad (3.16)$$

### Logarithms

We use the following notations:

$$\lg n = \log_2 n \quad (\text{binary logarithm}),$$

$$\ln n = \log_e n \quad (\text{natural logarithm}),$$

$$\lg^k n = (\lg n)^k \quad (\text{exponentiation}),$$

$$\lg \lg n = \lg(\lg n) \quad (\text{composition}).$$

We adopt the following notational convention: in the absence of parentheses, *a logarithm function applies only to the next term in the formula*, so that  $\lg n + 1$  means  $(\lg n) + 1$  and not  $\lg(n + 1)$ .

For any constant  $b > 1$ , the function  $\log_b n$  is undefined if  $n \leq 0$ , strictly increasing if  $n > 0$ , negative if  $0 < n < 1$ , positive if  $n > 1$ , and 0 if  $n = 1$ . For all real  $a > 0$ ,  $b > 0$ ,  $c > 0$ , and  $n$ , we have

$$a = b^{\log_b a}, \quad (3.17)$$

$$\log_c(ab) = \log_c a + \log_c b, \quad (3.18)$$

$$\log_b a^n = n \log_b a,$$

$$\log_b a = \frac{\log_c a}{\log_c b}, \quad (3.19)$$

$$\log_b(1/a) = -\log_b a, \quad (3.20)$$

$$\log_b a = \frac{1}{\log_a b},$$

$$a^{\log_b c} = c^{\log_b a}, \quad (3.21)$$

where, in each equation above, logarithm bases are not 1.

By equation (3.19), changing the base of a logarithm from one constant to another changes the value of the logarithm by only a constant factor. Consequently, we often use the notation “ $\lg n$ ” when we don’t care about constant factors, such as in  $O$ -notation. Computer scientists find 2 to be the most natural base for logarithms because so many algorithms and data structures involve splitting a problem into two parts.

There is a simple series expansion for  $\ln(1 + x)$  when  $|x| < 1$ :

$$\ln(1 + x) = x - \frac{x^2}{2} + \frac{x^3}{3} - \frac{x^4}{4} + \frac{x^5}{5} - \dots . \quad (3.22)$$

We also have the following inequalities for  $x > -1$ :

$$\frac{x}{1+x} \leq \ln(1+x) \leq x , \quad (3.23)$$

where equality holds only for  $x = 0$ .

We say that a function  $f(n)$  is *polylogarithmically bounded* if  $f(n) = O(\lg^k n)$  for some constant  $k$ . We can relate the growth of polynomials and polylogarithms by substituting  $\lg n$  for  $n$  and  $2^a$  for  $a$  in equation (3.13). For all real constants  $a > 0$  and  $b$ , we have

$$\lg^b n = o(n^a) . \quad (3.24)$$

Thus, any positive polynomial function grows faster than any polylogarithmic function.

### Factorials

The notation  $n!$  (read “ $n$  factorial”) is defined for integers  $n \geq 0$  as

$$n! = \begin{cases} 1 & \text{if } n = 0 , \\ n \cdot (n-1)! & \text{if } n > 0 . \end{cases}$$

Thus,  $n! = 1 \cdot 2 \cdot 3 \cdots n$ .

A weak upper bound on the factorial function is  $n! \leq n^n$ , since each of the  $n$  terms in the factorial product is at most  $n$ . *Stirling’s approximation*,

$$n! = \sqrt{2\pi n} \left(\frac{n}{e}\right)^n \left(1 + \Theta\left(\frac{1}{n}\right)\right) , \quad (3.25)$$

where  $e$  is the base of the natural logarithm, gives us a tighter upper bound, and a lower bound as well. Exercise 3.3-4 asks you to prove the three facts

$$n! = o(n^n) , \quad (3.26)$$

$$n! = \omega(2^n) , \quad (3.27)$$

$$\lg(n!) = \Theta(n \lg n) , \quad (3.28)$$

where Stirling's approximation is helpful in proving equation (3.28). The following equation also holds for all  $n \geq 1$ :

$$n! = \sqrt{2\pi n} \left(\frac{n}{e}\right)^n e^{\alpha_n} \quad (3.29)$$

where

$$\frac{1}{12n+1} < \alpha_n < \frac{1}{12n} .$$

### Functional iteration

We use the notation  $f^{(i)}(n)$  to denote the function  $f(n)$  iteratively applied  $i$  times to an initial value of  $n$ . Formally, let  $f(n)$  be a function over the reals. For non-negative integers  $i$ , we recursively define

$$f^{(i)}(n) = \begin{cases} n & \text{if } i = 0 , \\ f(f^{(i-1)}(n)) & \text{if } i > 0 . \end{cases} \quad (3.30)$$

For example, if  $f(n) = 2n$ , then  $f^{(i)}(n) = 2^i n$ .

### The iterated logarithm function

We use the notation  $\lg^* n$  (read “log star of  $n$ ”) to denote the iterated logarithm, defined as follows. Let  $\lg^{(i)} n$  be as defined above, with  $f(n) = \lg n$ . Because the logarithm of a nonpositive number is undefined,  $\lg^{(i)} n$  is defined only if  $\lg^{(i-1)} n > 0$ . Be sure to distinguish  $\lg^{(i)} n$  (the logarithm function applied  $i$  times in succession, starting with argument  $n$ ) from  $\lg^i n$  (the logarithm of  $n$  raised to the  $i$ th power). Then we define the iterated logarithm function as

$$\lg^* n = \min \{i \geq 0 : \lg^{(i)} n \leq 1\} .$$

The iterated logarithm is a *very* slowly growing function:

$$\begin{aligned} \lg^* 2 &= 1 , \\ \lg^* 4 &= 2 , \\ \lg^* 16 &= 3 , \\ \lg^* 65536 &= 4 , \\ \lg^*(2^{65536}) &= 5 . \end{aligned}$$

Since the number of atoms in the observable universe is estimated to be about  $10^{80}$ , which is much less than  $2^{65536} = 10^{65536/\lg 10} \approx 10^{19,728}$ , we rarely encounter an input size  $n$  for which  $\lg^* n > 5$ .

### Fibonacci numbers

We define the *Fibonacci numbers*  $F_i$ , for  $i \geq 0$ , as follows:

$$F_i = \begin{cases} 0 & \text{if } i = 0, \\ 1 & \text{if } i = 1, \\ F_{i-1} + F_{i-2} & \text{if } i \geq 2. \end{cases} \quad (3.31)$$

Thus, after the first two, each Fibonacci number is the sum of the two previous ones, yielding the sequence

0, 1, 1, 2, 3, 5, 8, 13, 21, 34, 55, ...

Fibonacci numbers are related to the *golden ratio*  $\phi$  and its conjugate  $\hat{\phi}$ , which are the two roots of the equation

$$x^2 = x + 1.$$

As Exercise 3.3-7 asks you to prove, the golden ratio is given by

$$\begin{aligned} \phi &= \frac{1 + \sqrt{5}}{2} \\ &= 1.61803\dots, \end{aligned} \quad (3.32)$$

and its conjugate, by

$$\begin{aligned} \hat{\phi} &= \frac{1 - \sqrt{5}}{2} \\ &= -.61803\dots \end{aligned} \quad (3.33)$$

Specifically, we have

$$F_i = \frac{\phi^i - \hat{\phi}^i}{\sqrt{5}},$$

which can be proved by induction (Exercise 3.3-8). Since  $|\hat{\phi}| < 1$ , we have

$$\begin{aligned} \frac{|\hat{\phi}^i|}{\sqrt{5}} &< \frac{1}{\sqrt{5}} \\ &< \frac{1}{2}, \end{aligned}$$

which implies that

$$F_i = \left\lfloor \frac{\phi^i}{\sqrt{5}} + \frac{1}{2} \right\rfloor, \quad (3.34)$$

which is to say that the  $i$ th Fibonacci number  $F_i$  is equal to  $\phi^i / \sqrt{5}$  rounded to the nearest integer. Thus, Fibonacci numbers grow exponentially.

**Exercises****3.3-1**

Show that if  $f(n)$  and  $g(n)$  are monotonically increasing functions, then so are the functions  $f(n) + g(n)$  and  $f(g(n))$ , and if  $f(n)$  and  $g(n)$  are in addition nonnegative, then  $f(n) \cdot g(n)$  is monotonically increasing.

**3.3-2**

Prove that  $\lfloor \alpha n \rfloor + \lceil (1 - \alpha)n \rceil = n$  for any integer  $n$  and real number  $\alpha$  in the range  $0 \leq \alpha \leq 1$ .

**3.3-3**

Use equation (3.14) or other means to show that  $(n + o(n))^k = \Theta(n^k)$  for any real constant  $k$ . Conclude that  $\lceil n \rceil^k = \Theta(n^k)$  and  $\lfloor n \rfloor^k = \Theta(n^k)$ .

**3.3-4**

Prove the following:

- a. Equation (3.21).
- b. Equations (3.26)–(3.28).
- c.  $\lg(\Theta(n)) = \Theta(\lg n)$ .

**★ 3.3-5**

Is the function  $\lceil \lg n \rceil!$  polynomially bounded? Is the function  $\lceil \lg \lg n \rceil!$  polynomially bounded?

**★ 3.3-6**

Which is asymptotically larger:  $\lg(\lg^* n)$  or  $\lg^*(\lg n)$ ?

**3.3-7**

Show that the golden ratio  $\phi$  and its conjugate  $\hat{\phi}$  both satisfy the equation  $x^2 = x + 1$ .

**3.3-8**

Prove by induction that the  $i$ th Fibonacci number satisfies the equation

$$F_i = (\phi^i - \hat{\phi}^i) / \sqrt{5},$$

where  $\phi$  is the golden ratio and  $\hat{\phi}$  is its conjugate.

**3.3-9**

Show that  $k \lg k = \Theta(n)$  implies  $k = \Theta(n / \lg n)$ .

---

**Problems**
**3-1 Asymptotic behavior of polynomials**

Let

$$p(n) = \sum_{i=0}^d a_i n^i,$$

where  $a_d > 0$ , be a degree- $d$  polynomial in  $n$ , and let  $k$  be a constant. Use the definitions of the asymptotic notations to prove the following properties.

- a. If  $k \geq d$ , then  $p(n) = O(n^k)$ .
- b. If  $k \leq d$ , then  $p(n) = \Omega(n^k)$ .
- c. If  $k = d$ , then  $p(n) = \Theta(n^k)$ .
- d. If  $k > d$ , then  $p(n) = o(n^k)$ .
- e. If  $k < d$ , then  $p(n) = \omega(n^k)$ .

**3-2 Relative asymptotic growths**

Indicate, for each pair of expressions  $(A, B)$  in the table below whether  $A$  is  $O$ ,  $o$ ,  $\Omega$ ,  $\omega$ , or  $\Theta$  of  $B$ . Assume that  $k \geq 1$ ,  $\epsilon > 0$ , and  $c > 1$  are constants. Write your answer in the form of the table with “yes” or “no” written in each box.

	$A$	$B$	$O$	$o$	$\Omega$	$\omega$	$\Theta$
a.	$\lg^k n$	$n^\epsilon$					
b.	$n^k$	$c^n$					
c.	$\sqrt{n}$	$n^{\sin n}$					
d.	$2^n$	$2^{n/2}$					
e.	$n^{\lg c}$	$c^{\lg n}$					
f.	$\lg(n!)$	$\lg(n^n)$					

**3-3 Ordering by asymptotic growth rates**

- a. Rank the following functions by order of growth. That is, find an arrangement  $g_1, g_2, \dots, g_{30}$  of the functions satisfying  $g_1 = \Omega(g_2)$ ,  $g_2 = \Omega(g_3)$ ,  $\dots$ ,  $g_{29} = \Omega(g_{30})$ . Partition your list into equivalence classes such that functions  $f(n)$  and  $g(n)$  belong to the same class if and only if  $f(n) = \Theta(g(n))$ .

$\lg(\lg^* n)$	$2^{\lg^* n}$	$(\sqrt{2})^{\lg n}$	$n^2$	$n!$	$(\lg n)!$
$(3/2)^n$	$n^3$	$\lg^2 n$	$\lg(n!)$	$2^{2^n}$	$n^{1/\lg n}$
$\ln \ln n$	$\lg^* n$	$n \cdot 2^n$	$n^{\lg \lg n}$	$\ln n$	1
$2^{\lg n}$	$(\lg n)^{\lg n}$	$e^n$	$4^{\lg n}$	$(n+1)!$	$\sqrt{\lg n}$
$\lg^*(\lg n)$	$2^{\sqrt{2 \lg n}}$	$n$	$2^n$	$n \lg n$	$2^{2^{n+1}}$

- b.** Give an example of a single nonnegative function  $f(n)$  such that for all functions  $g_i(n)$  in part (a),  $f(n)$  is neither  $O(g_i(n))$  nor  $\Omega(g_i(n))$ .

### 3-4 Asymptotic notation properties

Let  $f(n)$  and  $g(n)$  be asymptotically positive functions. Prove or disprove each of the following conjectures.

- a.**  $f(n) = O(g(n))$  implies  $g(n) = O(f(n))$ .
- b.**  $f(n) + g(n) = \Theta(\min\{f(n), g(n)\})$ .
- c.**  $f(n) = O(g(n))$  implies  $\lg f(n) = O(\lg g(n))$ , where  $\lg g(n) \geq 1$  and  $f(n) \geq 1$  for all sufficiently large  $n$ .
- d.**  $f(n) = O(g(n))$  implies  $2^{f(n)} = O(2^{g(n)})$ .
- e.**  $f(n) = O((f(n))^2)$ .
- f.**  $f(n) = O(g(n))$  implies  $g(n) = \Omega(f(n))$ .
- g.**  $f(n) = \Theta(f(n/2))$ .
- h.**  $f(n) + o(f(n)) = \Theta(f(n))$ .

### 3-5 Manipulating asymptotic notation

Let  $f(n)$  and  $g(n)$  be asymptotically positive functions. Prove the following identities:

- a.**  $\Theta(\Theta(f(n))) = \Theta(f(n))$ .
- b.**  $\Theta(f(n)) + O(f(n)) = \Theta(f(n))$ .
- c.**  $\Theta(f(n)) + \Theta(g(n)) = \Theta(f(n) + g(n))$ .
- d.**  $\Theta(f(n)) \cdot \Theta(g(n)) = \Theta(f(n) \cdot g(n))$ .



- e. Argue that for any real constants  $a_1, b_1 > 0$  and integer constants  $k_1, k_2$ , the following asymptotic bound holds:

$$(a_1 n)^{k_1} \lg^{k_2}(a_2 n) = \Theta(n^{k_1} \lg^{k_2} n) .$$

- ★ f. Prove that for  $S \subseteq \mathbb{Z}$ , we have

$$\sum_{k \in S} \Theta(f(k)) = \Theta \left( \sum_{k \in S} f(k) \right) ,$$

assuming that both sums converge.

- ★ g. Show that for  $S \subseteq \mathbb{Z}$ , the following asymptotic bound does not necessarily hold, even assuming that both products converge, by giving a counterexample:

$$\prod_{k \in S} \Theta(f(k)) = \Theta \left( \prod_{k \in S} f(k) \right) .$$

### 3-6 Variations on $O$ and $\Omega$

Some authors define  $\Omega$ -notation in a slightly different way than this textbook does. We'll use the nomenclature  $\tilde{\Omega}$  (read “omega infinity”) for this alternative definition. We say that  $f(n) = \tilde{\Omega}(g(n))$  if there exists a positive constant  $c$  such that  $f(n) \geq c g(n) \geq 0$  for infinitely many integers  $n$ .

- Show that for any two asymptotically nonnegative functions  $f(n)$  and  $g(n)$ , we have  $f(n) = O(g(n))$  or  $f(n) = \tilde{\Omega}(g(n))$  (or both).
- Show that there exist two asymptotically nonnegative functions  $f(n)$  and  $g(n)$  for which neither  $f(n) = O(g(n))$  nor  $f(n) = \tilde{\Omega}(g(n))$  holds.
- Describe the potential advantages and disadvantages of using  $\tilde{\Omega}$ -notation instead of  $\Omega$ -notation to characterize the running times of programs.

Some authors also define  $O$  in a slightly different manner. We'll use  $O'$  for the alternative definition:  $f(n) = O'(g(n))$  if and only if  $|f(n)| = O(g(n))$ .

- What happens to each direction of the “if and only if” in Theorem 3.1 on page 56 if we substitute  $O'$  for  $O$  but still use  $\Omega$ ?

Some authors define  $\tilde{O}$  (read “soft-oh”) to mean  $O$  with logarithmic factors ignored:

$$\widetilde{O}(g(n)) = \{f(n) : \text{there exist positive constants } c, k, \text{ and } n_0 \text{ such that} \\ 0 \leq f(n) \leq cg(n) \lg^k(n) \text{ for all } n \geq n_0\}.$$

- e. Define  $\widetilde{\Omega}$  and  $\widetilde{\Theta}$  in a similar manner. Prove the corresponding analog to Theorem 3.1.

### 3-7 Iterated functions

We can apply the iteration operator  $*$  used in the  $\lg^*$  function to any monotonically increasing function  $f(n)$  over the reals. For a given constant  $c \in \mathbb{R}$ , we define the iterated function  $f_c^*$  by

$$f_c^*(n) = \min \{i \geq 0 : f^{(i)}(n) \leq c\},$$

which need not be well defined in all cases. In other words, the quantity  $f_c^*(n)$  is the minimum number of iterated applications of the function  $f$  required to reduce its argument down to  $c$  or less.

For each of the functions  $f(n)$  and constants  $c$  in the table below, give as tight a bound as possible on  $f_c^*(n)$ . If there is no  $i$  such that  $f^{(i)}(n) \leq c$ , write “undefined” as your answer.

	$f(n)$	$c$	$f_c^*(n)$
a.	$n - 1$	0	
b.	$\lg n$	1	
c.	$n/2$	1	
d.	$n/2$	2	
e.	$\sqrt{n}$	2	
f.	$\sqrt{n}$	1	
g.	$n^{1/3}$	2	

---

## Chapter notes

Knuth [259] traces the origin of the  $O$ -notation to a number-theory text by P. Bachmann in 1892. The  $o$ -notation was invented by E. Landau in 1909 for his discussion of the distribution of prime numbers. The  $\Omega$  and  $\Theta$  notations were advocated by Knuth [265] to correct the popular, but technically sloppy, practice in the literature of using  $O$ -notation for both upper and lower bounds. As noted earlier in this chapter, many people continue to use the  $O$ -notation where the  $\Theta$ -notation is more technically precise. The soft-oh notation  $\widetilde{O}$  in Problem 3-6 was introduced

by Babai, Luks, and Seress [31], although it was originally written as  $O\sim$ . Some authors now define  $\widetilde{O}(g(n))$  as ignoring factors that are logarithmic in  $g(n)$ , rather than in  $n$ . With this definition, we can say that  $n2^n = \widetilde{O}(2^n)$ , but with the definition in Problem 3-6, this statement is not true. Further discussion of the history and development of asymptotic notations appears in works by Knuth [259, 265] and Brassard and Bratley [70].

Not all authors define the asymptotic notations in the same way, although the various definitions agree in most common situations. Some of the alternative definitions encompass functions that are not asymptotically nonnegative, as long as their absolute values are appropriately bounded.

Equation (3.29) is due to Robbins [381]. Other properties of elementary mathematical functions can be found in any good mathematical reference, such as Abramowitz and Stegun [1] or Zwillinger [468], or in a calculus book, such as Apostol [19] or Thomas et al. [433]. Knuth [259] and Graham, Knuth, and Patashnik [199] contain a wealth of material on discrete mathematics as used in computer science.

---

## 11 Hash Tables

Many applications require a dynamic set that supports only the dictionary operations INSERT, SEARCH, and DELETE. For example, a compiler that translates a programming language maintains a symbol table, in which the keys of elements are arbitrary character strings corresponding to identifiers in the language. A hash table is an effective data structure for implementing dictionaries. Although searching for an element in a hash table can take as long as searching for an element in a linked list— $\Theta(n)$  time in the worst case—in practice, hashing performs extremely well. Under reasonable assumptions, the average time to search for an element in a hash table is  $O(1)$ . Indeed, the built-in dictionaries of Python are implemented with hash tables.

A hash table generalizes the simpler notion of an ordinary array. Directly addressing into an ordinary array takes advantage of the  $O(1)$  access time for any array element. Section 11.1 discusses direct addressing in more detail. To use direct addressing, you must be able to allocate an array that contains a position for every possible key.

When the number of keys actually stored is small relative to the total number of possible keys, hash tables become an effective alternative to directly addressing an array, since a hash table typically uses an array of size proportional to the number of keys actually stored. Instead of using the key as an array index directly, we *compute* the array index from the key. Section 11.2 presents the main ideas, focusing on “chaining” as a way to handle “collisions,” in which more than one key maps to the same array index. Section 11.3 describes how to compute array indices from keys using hash functions. We present and analyze several variations on the basic theme. Section 11.4 looks at “open addressing,” which is another way to deal with collisions. The bottom line is that hashing is an extremely effective and practical technique: the basic dictionary operations require only  $O(1)$  time on the average. Section 11.5 discusses the hierarchical memory systems of modern computer systems have and illustrates how to design hash tables that work well in such systems.


## 11.1 Direct-address tables

Direct addressing is a simple technique that works well when the universe  $U$  of keys is reasonably small. Suppose that an application needs a dynamic set in which each element has a distinct key drawn from the universe  $U = \{0, 1, \dots, m-1\}$ , where  $m$  is not too large.

To represent the dynamic set, you can use an array, or *direct-address table*, denoted by  $T[0 : m-1]$ , in which each position, or *slot*, corresponds to a key in the universe  $U$ . Figure 11.1 illustrates this approach. Slot  $k$  points to an element in the set with key  $k$ . If the set contains no element with key  $k$ , then  $T[k] = \text{NIL}$ .

The dictionary operations DIRECT-ADDRESS-SEARCH, DIRECT-ADDRESS-INSERT, and DIRECT-ADDRESS-DELETE on the following page are trivial to implement. Each takes only  $O(1)$  time.

For some applications, the direct-address table itself can hold the elements in the dynamic set. That is, rather than storing an element's key and satellite data in an object external to the direct-address table, with a pointer from a slot in the table to the object, save space by storing the object directly in the slot. To indicate an empty slot, use a special key. Then again, why store the key of the object at all? The index of the object *is* its key! Of course, then you'd need some way to tell whether slots are empty.



**Figure 11.1** How to implement a dynamic set by a direct-address table  $T$ . Each key in the universe  $U = \{0, 1, \dots, 9\}$  corresponds to an index into the table. The set  $K = \{2, 3, 5, 8\}$  of actual keys determines the slots in the table that contain pointers to elements. The other slots, in blue, contain NIL.

```
DIRECT-ADDRESS-SEARCH( $T, k$ )
```

```
1  return  $T[k]$ 
```

```
DIRECT-ADDRESS-INSERT( $T, x$ )
```

```
1   $T[x.key] = x$ 
```

```
DIRECT-ADDRESS-DELETE( $T, x$ )
```

```
1   $T[x.key] = \text{NIL}$ 
```

## Exercises

### 11.1-1

A dynamic set  $S$  is represented by a direct-address table  $T$  of length  $m$ . Describe a procedure that finds the maximum element of  $S$ . What is the worst-case performance of your procedure?

### 11.1-2

A *bit vector* is simply an array of bits (each either 0 or 1). A bit vector of length  $m$  takes much less space than an array of  $m$  pointers. Describe how to use a bit vector to represent a dynamic set of distinct elements drawn from the set  $\{0, 1, \dots, m-1\}$  and with no satellite data. Dictionary operations should run in  $O(1)$  time.

### 11.1-3

Suggest how to implement a direct-address table in which the keys of stored elements do not need to be distinct and the elements can have satellite data. All three dictionary operations (INSERT, DELETE, and SEARCH) should run in  $O(1)$  time. (Don't forget that DELETE takes as an argument a pointer to an object to be deleted, not a key.)

### ★ 11.1-4

Suppose that you want to implement a dictionary by using direct addressing on a *huge* array. That is, if the array size is  $m$  and the dictionary contains at most  $n$  elements at any one time, then  $m \gg n$ . At the start, the array entries may contain garbage, and initializing the entire array is impractical because of its size. Describe a scheme for implementing a direct-address dictionary on a huge array. Each stored object should use  $O(1)$  space; the operations SEARCH, INSERT, and DELETE should take  $O(1)$  time each; and initializing the data structure should take  $O(1)$  time. (*Hint:* Use an additional array, treated somewhat like a stack whose size is the number of keys actually stored in the dictionary, to help determine whether a given entry in the huge array is valid or not.)

---

## 11.2 Hash tables

The downside of direct addressing is apparent: if the universe  $U$  is large or infinite, storing a table  $T$  of size  $|U|$  may be impractical, or even impossible, given the memory available on a typical computer. Furthermore, the set  $K$  of keys *actually stored* may be so small relative to  $U$  that most of the space allocated for  $T$  would be wasted.

When the set  $K$  of keys stored in a dictionary is much smaller than the universe  $U$  of all possible keys, a hash table requires much less storage than a direct-address table. Specifically, the storage requirement reduces to  $\Theta(|K|)$  while maintaining the benefit that searching for an element in the hash table still requires only  $O(1)$  time. The catch is that this bound is for the *average-case time*,<sup>1</sup> whereas for direct addressing it holds for the *worst-case time*.

With direct addressing, an element with key  $k$  is stored in slot  $k$ , but with hashing, we use a **hash function**  $h$  to compute the slot number from the key  $k$ , so that the element goes into slot  $h(k)$ . The hash function  $h$  maps the universe  $U$  of keys into the slots of a **hash table**  $T[0 : m - 1]$ :

$$h : U \rightarrow \{0, 1, \dots, m - 1\} ,$$


where the size  $m$  of the hash table is typically much less than  $|U|$ . We say that an element with key  $k$  **hashes** to slot  $h(k)$ , and we also say that  $h(k)$  is the **hash value** of key  $k$ . Figure 11.2 illustrates the basic idea. The hash function reduces the range of array indices and hence the size of the array. Instead of a size of  $|U|$ , the array can have size  $m$ . An example of a simple, but not particularly good, hash function is  $h(k) = k \bmod m$ .

There is one hitch, namely that two keys may hash to the same slot. We call this situation a **collision**. Fortunately, there are effective techniques for resolving the conflict created by collisions.

Of course, the ideal solution is to avoid collisions altogether. We might try to achieve this goal by choosing a suitable hash function  $h$ . One idea is to make  $h$  appear to be “random,” thus avoiding collisions or at least minimizing their number. The very term “to hash,” evoking images of random mixing and chopping, captures the spirit of this approach. (Of course, a hash function  $h$  must be deterministic in that a given input  $k$  must always produce the same output  $h(k)$ .) Because  $|U| > m$ , however, there must be at least two keys that have the same hash value,

---

<sup>1</sup> The definition of “average-case” requires care—are we assuming an input distribution over the keys, or are we randomizing the choice of hash function itself? We’ll consider both approaches, but with an emphasis on the use of a randomly chosen hash function.



**Figure 11.2** Using a hash function  $h$  to map keys to hash-table slots. Because keys  $k_2$  and  $k_5$  map to the same slot, they collide.

and avoiding collisions altogether is impossible. Thus, although a well-designed, “random”-looking hash function can reduce the number of collisions, we still need a method for resolving the collisions that do occur.

The remainder of this section first presents a definition of “independent uniform hashing,” which captures the simplest notion of what it means for a hash function to be “random.” It then presents and analyzes the simplest collision resolution technique, called chaining. Section 11.4 introduces an alternative method for resolving collisions, called open addressing.


### Independent uniform hashing

An “ideal” hashing function  $h$  would have, for each possible input  $k$  in the domain  $U$ , an output  $h(k)$  that is an element randomly and independently chosen uniformly from the range  $\{0, 1, \dots, m-1\}$ . Once a value  $h(k)$  is randomly chosen, each subsequent call to  $h$  with the same input  $k$  yields the same output  $h(k)$ .

We call such an ideal hash function an *independent uniform hash function*. Such a function is also often called a *random oracle* [43]. When hash tables are implemented with an independent uniform hash function, we say we are using *independent uniform hashing*.

Independent uniform hashing is an ideal theoretical abstraction, but it is not something that can reasonably be implemented in practice. Nonetheless, we’ll analyze the efficiency of hashing under the assumption of independent uniform hashing and then present ways of achieving useful practical approximations to this ideal.





**Figure 11.3** Collision resolution by chaining. Each nonempty hash-table slot  $T[j]$  points to a linked list of all the keys whose hash value is  $j$ . For example,  $h(k_1) = h(k_4)$  and  $h(k_5) = h(k_2) = h(k_7)$ . The list can be either singly or doubly linked. We show it as doubly linked because deletion may be faster that way when the deletion procedure knows which list element (not just which key) is to be deleted.

### Collision resolution by chaining

At a high level, you can think of hashing with chaining as a nonrecursive form of divide-and-conquer: the input set of  $n$  elements is divided randomly into  $m$  subsets, each of approximate size  $n/m$ . A hash function determines which subset an element belongs to. Each subset is managed independently as a list.

Figure 11.3 shows the idea behind **chaining**: each nonempty slot points to a linked list, and all the elements that hash to the same slot go into that slot's linked list. Slot  $j$  contains a pointer to the head of the list of all stored elements with hash value  $j$ . If there are no such elements, then slot  $j$  contains NIL.

When collisions are resolved by chaining, the dictionary operations are straightforward to implement. They appear on the next page and use the linked-list procedures from Section 10.2. The worst-case running time for insertion is  $O(1)$ . The insertion procedure is fast in part because it assumes that the element  $x$  being inserted is not already present in the table. To enforce this assumption, you can search (at additional cost) for an element whose key is  $x.key$  before inserting. For searching, the worst-case running time is proportional to the length of the list. (We'll analyze this operation more closely below.) Deletion takes  $O(1)$  time if the lists are doubly linked, as in Figure 11.3. (Since CHAINED-HASH-DELETE takes as input an element  $x$  and not its key  $k$ , no search is needed. If the hash table supports deletion, then its linked lists should be doubly linked in order to delete an item quickly. If the lists were only singly linked, then by Exercise 10.2-1, deletion

```

CHAINED-HASH-INSERT( $T, x$ )
1  LIST-PREPEND( $T[h(x.key)], x$ )

CHAINED-HASH-SEARCH( $T, k$ )
1  return LIST-SEARCH( $T[h(k)], k$ )

CHAINED-HASH-DELETE( $T, x$ )
1  LIST-DELETE( $T[h(x.key)], x$ )

```

could take time proportional to the length of the list. With singly linked lists, both deletion and searching would have the same asymptotic running times.)

### Analysis of hashing with chaining

How well does hashing with chaining perform? In particular, how long does it take to search for an element with a given key?

Given a hash table  $T$  with  $m$  slots that stores  $n$  elements, we define the *load factor*  $\alpha$  for  $T$  as  $n/m$ , that is, the average number of elements stored in a chain. Our analysis will be in terms of  $\alpha$ , which can be less than, equal to, or greater than 1.

The worst-case behavior of hashing with chaining is terrible: all  $n$  keys hash to the same slot, creating a list of length  $n$ . The worst-case time for searching is thus  $\Theta(n)$  plus the time to compute the hash function—no better than using one linked list for all the elements. We clearly don't use hash tables for their worst-case performance.

The average-case performance of hashing depends on how well the hash function  $h$  distributes the set of keys to be stored among the  $m$  slots, on the average (meaning with respect to the distribution of keys to be hashed and with respect to the choice of hash function, if this choice is randomized). Section 11.3 discusses these issues, but for now we assume that any given element is equally likely to hash into any of the  $m$  slots. That is, the hash function is *uniform*. We further assume that where a given element hashes to is *independent* of where any other elements hash to. In other words, we assume that we are using *independent uniform hashing*.

Because hashes of distinct keys are assumed to be independent, independent uniform hashing is *universal*: the chance that any two distinct keys  $k_1$  and  $k_2$  collide is at most  $1/m$ . Universality is important in our analysis and also in the specification of universal families of hash functions, which we'll see in Section 11.3.2.

For  $j = 0, 1, \dots, m - 1$ , denote the length of the list  $T[j]$  by  $n_j$ , so that

$$n = n_0 + n_1 + \cdots + n_{m-1}, \quad (11.1)$$

and the expected value of  $n_j$  is  $E[n_j] = \alpha = n/m$ .

We assume that  $O(1)$  time suffices to compute the hash value  $h(k)$ , so that the time required to search for an element with key  $k$  depends linearly on the length  $n_{h(k)}$  of the list  $T[h(k)]$ . Setting aside the  $O(1)$  time required to compute the hash function and to access slot  $h(k)$ , we'll consider the expected number of elements examined by the search algorithm, that is, the number of elements in the list  $T[h(k)]$  that the algorithm checks to see whether any have a key equal to  $k$ . We consider two cases. In the first, the search is unsuccessful: no element in the table has key  $k$ . In the second, the search successfully finds an element with key  $k$ .

**Theorem 11.1**

In a hash table in which collisions are resolved by chaining, an unsuccessful search takes  $\Theta(1 + \alpha)$  time on average, under the assumption of independent uniform hashing.

**Proof** Under the assumption of independent uniform hashing, any key  $k$  not already stored in the table is equally likely to hash to any of the  $m$  slots. The expected time to search unsuccessfully for a key  $k$  is the expected time to search to the end of list  $T[h(k)]$ , which has expected length  $E[n_{h(k)}] = \alpha$ . Thus, the expected number of elements examined in an unsuccessful search is  $\alpha$ , and the total time required (including the time for computing  $h(k)$ ) is  $\Theta(1 + \alpha)$ . ■

The situation for a successful search is slightly different. An unsuccessful search is equally likely to go to any slot of the hash table. A successful search, however, cannot go to an empty slot, since it is for an element that is present in one of the linked lists. We assume that the element searched for is equally likely to be any one of the elements in the table, so the longer the list, the more likely that the search is for one of its elements. Even so, the expected search time still turns out to be  $\Theta(1 + \alpha)$ .

**Theorem 11.2**

In a hash table in which collisions are resolved by chaining, a successful search takes  $\Theta(1 + \alpha)$  time on average, under the assumption of independent uniform hashing.

**Proof** We assume that the element being searched for is equally likely to be any of the  $n$  elements stored in the table. The number of elements examined during a successful search for an element  $x$  is 1 more than the number of elements that appear before  $x$  in  $x$ 's list. Because new elements are placed at the front of the list,

elements before  $x$  in the list were all inserted after  $x$  was inserted. Let  $x_i$  denote the  $i$ th element inserted into the table, for  $i = 1, 2, \dots, n$ , and let  $k_i = x_i.\text{key}$ .

Our analysis uses indicator random variables extensively. For each slot  $q$  in the table and for each pair of distinct keys  $k_i$  and  $k_j$ , we define the indicator random variable

$$X_{ijq} = \mathbf{I} \{ \text{the search is for } x_i, h(k_i) = q, \text{ and } h(k_j) = q \} .$$

That is,  $X_{ijq} = 1$  when keys  $k_i$  and  $k_j$  collide at slot  $q$  and the search is for element  $x_i$ . Because  $\Pr \{ \text{the search is for } x_i \} = 1/n$ ,  $\Pr \{ h(k_i) = q \} = 1/m$ ,  $\Pr \{ h(k_j) = q \} = 1/m$ , and these events are all independent, we have that  $\Pr \{ X_{ijq} = 1 \} = 1/nm^2$ . Lemma 5.1 on page 130 gives  $E[X_{ijq}] = 1/nm^2$ .

Next, we define, for each element  $x_j$ , the indicator random variable

$$\begin{aligned} Y_j &= \mathbf{I} \{ x_j \text{ appears in a list prior to the element being searched for} \} \\ &= \sum_{q=0}^{m-1} \sum_{i=1}^{j-1} X_{ijq} , \end{aligned}$$

since at most one of the  $X_{ijq}$  equals 1, namely when the element  $x_i$  being searched for belongs to the same list as  $x_j$  (pointed to by slot  $q$ ), and  $i < j$  (so that  $x_i$  appears after  $x_j$  in the list).

Our final random variable is  $Z$ , which counts how many elements appear in the list prior to the element being searched for:

$$Z = \sum_{j=1}^n Y_j .$$

Because we must count the element being searched for as well as all those preceding it in its list, we wish to compute  $E[Z + 1]$ . Using linearity of expectation (equation (C.24) on page 1192), we have

$$\begin{aligned} E[Z + 1] &= E \left[ 1 + \sum_{j=1}^n Y_j \right] \\ &= 1 + E \left[ \sum_{j=1}^n \sum_{q=0}^{m-1} \sum_{i=1}^{j-1} X_{ijq} \right] \\ &= 1 + E \left[ \sum_{q=0}^{m-1} \sum_{j=1}^n \sum_{i=1}^{j-1} X_{ijq} \right] \\ &= 1 + \sum_{q=0}^{m-1} \sum_{j=1}^n \sum_{i=1}^{j-1} E[X_{ijq}] \quad (\text{by linearity of expectation}) \end{aligned}$$

$$\begin{aligned}
&= 1 + \sum_{q=0}^{m-1} \sum_{j=1}^n \sum_{i=1}^{j-1} \frac{1}{nm^2} \\
&= 1 + m \cdot \frac{n(n-1)}{2} \cdot \frac{1}{nm^2} \quad (\text{by equation (A.2) on page 1141}) \\
&= 1 + \frac{n-1}{2m} \\
&= 1 + \frac{n}{2m} - \frac{1}{2m} \\
&= 1 + \frac{\alpha}{2} - \frac{\alpha}{2n}.
\end{aligned}$$

Thus, the total time required for a successful search (including the time for computing the hash function) is  $\Theta(2 + \alpha/2 - \alpha/2n) = \Theta(1 + \alpha)$ . ■

What does this analysis mean? If the number of elements in the table is at most proportional to the number of hash-table slots, we have  $n = O(m)$  and, consequently,  $\alpha = n/m = O(m)/m = O(1)$ . Thus, searching takes constant time on average. Since insertion takes  $O(1)$  worst-case time and deletion takes  $O(1)$  worst-case time when the lists are doubly linked (assuming that the list element to be deleted is known, and not just its key), we can support all dictionary operations in  $O(1)$  time on average.

The analysis in the preceding two theorems depends only on two essential properties of independent uniform hashing: uniformity (each key is equally likely to hash to any one of the  $m$  slots), and independence (so any two distinct keys collide with probability  $1/m$ ).

## Exercises

### 11.2-1

You use a hash function  $h$  to hash  $n$  distinct keys into an array  $T$  of length  $m$ . Assuming independent uniform hashing, what is the expected number of collisions? More precisely, what is the expected cardinality of  $\{\{k_1, k_2\} : k_1 \neq k_2 \text{ and } h(k_1) = h(k_2)\}$ ?

### 11.2-2

Consider a hash table with 9 slots and the hash function  $h(k) = k \bmod 9$ . Demonstrate what happens upon inserting the keys 5, 28, 19, 15, 20, 33, 12, 17, 10 with collisions resolved by chaining.

**11.2-3**

Professor Marley hypothesizes that he can obtain substantial performance gains by modifying the chaining scheme to keep each list in sorted order. How does the professor's modification affect the running time for successful searches, unsuccessful searches, insertions, and deletions?

**11.2-4**

Suggest how to allocate and deallocate storage for elements within the hash table itself by creating a “free list”: a linked list of all the unused slots. Assume that one slot can store a flag and either one element plus a pointer or two pointers. All dictionary and free-list operations should run in  $O(1)$  expected time. Does the free list need to be doubly linked, or does a singly linked free list suffice?

**11.2-5**

You need to store a set of  $n$  keys in a hash table of size  $m$ . Show that if the keys are drawn from a universe  $U$  with  $|U| > (n - 1)m$ , then  $U$  has a subset of size  $n$  consisting of keys that all hash to the same slot, so that the worst-case searching time for hashing with chaining is  $\Theta(n)$ .

**11.2-6**

You have stored  $n$  keys in a hash table of size  $m$ , with collisions resolved by chaining, and you know the length of each chain, including the length  $L$  of the longest chain. Describe a procedure that selects a key uniformly at random from among the keys in the hash table and returns it in expected time  $O(L \cdot (1 + 1/\alpha))$ .

---

**11.3 Hash functions**

For hashing to work well, it needs a good hash function. Along with being efficiently computable, what properties does a good hash function have? How do you design good hash functions?

This section first attempts to answer these questions based on two ad hoc approaches for creating hash functions: hashing by division and hashing by multiplication. Although these methods work well for some sets of input keys, they are limited because they try to provide a single fixed hash function that works well on any data—an approach called *static hashing*.

We then see that provably good average-case performance for *any* data can be obtained by designing a suitable *family* of hash functions and choosing a hash function at random from this family at runtime, independent of the data to be hashed. The approach we examine is called random hashing. A particular kind of random

hashing, universal hashing, works well. As we saw with quicksort in Chapter 7, randomization is a powerful algorithmic design tool.

### What makes a good hash function?

A good hash function satisfies (approximately) the assumption of independent uniform hashing: each key is equally likely to hash to any of the  $m$  slots, independently of where any other keys have hashed to. What does “equally likely” mean here? If the hash function is fixed, any probabilities would have to be based on the probability distribution of the input keys.

Unfortunately, you typically have no way to check this condition, unless you happen to know the probability distribution from which the keys are drawn. Moreover, the keys might not be drawn independently.

Occasionally you might know the distribution. For example, if you know that the keys are random real numbers  $k$  independently and uniformly distributed in the range  $0 \leq k < 1$ , then the hash function

$$h(k) = \lfloor km \rfloor$$

satisfies the condition of independent uniform hashing.

A good static hashing approach derives the hash value in a way that you expect to be independent of any patterns that might exist in the data. For example, the “division method” (discussed in Section 11.3.1) computes the hash value as the remainder when the key is divided by a specified prime number. This method may give good results, if you (somehow) choose a prime number that is unrelated to any patterns in the distribution of keys.

Random hashing, described in Section 11.3.2, picks the hash function to be used at random from a suitable family of hashing functions. This approach removes any need to know anything about the probability distribution of the input keys, as the randomization necessary for good average-case behavior then comes from the (known) random process used to pick the hash function from the family of hash functions, rather than from the (unknown) process used to create the input keys. We recommend that you use random hashing.

### Keys are integers, vectors, or strings

In practice, a hash function is designed to handle keys that are one of the following two types:

- A short nonnegative integer that fits in a  $w$ -bit machine word. Typical values for  $w$  would be 32 or 64.

- A short vector of nonnegative integers, each of bounded size. For example, each element might be an 8-bit byte, in which case the vector is often called a (byte) string. The vector might be of variable length.

To begin, we assume that keys are short nonnegative integers. Handling vector keys is more complicated and discussed in Sections 11.3.5 and 11.5.2.

### 11.3.1 Static hashing

Static hashing uses a single, fixed hash function. The only randomization available is through the (usually unknown) distribution of input keys. This section discusses two standard approaches for static hashing: the division method and the multiplication method. Although static hashing is no longer recommended, the multiplication method also provides a good foundation for “nonstatic” hashing—better known as random hashing—where the hash function is chosen at random from a suitable family of hash functions.

#### The division method

The *division method* for creating hash functions maps a key  $k$  into one of  $m$  slots by taking the remainder of  $k$  divided by  $m$ . That is, the hash function is

$$h(k) = k \bmod m .$$

For example, if the hash table has size  $m = 12$  and the key is  $k = 100$ , then  $h(k) = 4$ . Since it requires only a single division operation, hashing by division is quite fast.

The division method may work well when  $m$  is a prime not too close to an exact power of 2. There is no guarantee that this method provides good average-case performance, however, and it may complicate applications since it constrains the size of the hash tables to be prime.


#### The multiplication method

The general *multiplication method* for creating hash functions operates in two steps. First, multiply the key  $k$  by a constant  $A$  in the range  $0 < A < 1$  and extract the fractional part of  $kA$ . Then, multiply this value by  $m$  and take the floor of the result. That is, the hash function is

$$h(k) = \lfloor m (kA \bmod 1) \rfloor ,$$

where “ $kA \bmod 1$ ” means the fractional part of  $kA$ , that is,  $kA - \lfloor kA \rfloor$ . The general multiplication method has the advantage that the value of  $m$  is not critical and you can choose it independently of how you choose the multiplicative constant  $A$ .





**Figure 11.4** The multiply-shift method to compute a hash function. The  $w$ -bit representation of the key  $k$  is multiplied by the  $w$ -bit value  $a = A \cdot 2^w$ . The  $\ell$  highest-order bits of the lower  $w$ -bit half of the product form the desired hash value  $h_a(k)$ .

### The multiply-shift method

In practice, the multiplication method is best in the special case where the number  $m$  of hash-table slots is an exact power of 2, so that  $m = 2^\ell$  for some integer  $\ell$ , where  $\ell \leq w$  and  $w$  is the number of bits in a machine word. If you choose a fixed  $w$ -bit positive integer  $a = A 2^w$ , where  $0 < A < 1$  as in the multiplication method so that  $a$  is in the range  $0 < a < 2^w$ , you can implement the function on most computers as follows. We assume that a key  $k$  fits into a single  $w$ -bit word.

Referring to Figure 11.4, first multiply  $k$  by the  $w$ -bit integer  $a$ . The result is a  $2w$ -bit value  $r_1 2^w + r_0$ , where  $r_1$  is the high-order  $w$ -bit word of the product and  $r_0$  is the low-order  $w$ -bit word of the product. The desired  $\ell$ -bit hash value consists of the  $\ell$  most significant bits of  $r_0$ . (Since  $r_1$  is ignored, the hash function can be implemented on a computer that produces only a  $w$ -bit product given two  $w$ -bit inputs, that is, where the multiplication operation computes modulo  $2^w$ .)

In other words, you define the hash function  $h = h_a$ , where

$$h_a(k) = (ka \bmod 2^w) \ggg (w - \ell) \quad (11.2)$$

for a fixed nonzero  $w$ -bit value  $a$ . Since the product  $ka$  of two  $w$ -bit words occupies  $2w$  bits, taking this product modulo  $2^w$  zeroes out the high-order  $w$  bits ( $r_1$ ), leaving only the low-order  $w$  bits ( $r_0$ ). The  $\ggg$  operator performs a logical right shift by  $w - \ell$  bits, shifting zeros into the vacated positions on the left, so that the  $\ell$  most significant bits of  $r_0$  move into the  $\ell$  rightmost positions. (It's the same as dividing by  $2^{w-\ell}$  and taking the floor of the result.) The resulting value equals the  $\ell$  most significant bits of  $r_0$ . The hash function  $h_a$  can be implemented with three machine instructions: multiplication, subtraction, and logical right shift.

As an example, suppose that  $k = 123456$ ,  $\ell = 14$ ,  $m = 2^{14} = 16384$ , and  $w = 32$ . Suppose further that we choose  $a = 2654435769$  (following a suggestion

of Knuth [261]). Then  $ka = 327706022297664 = (76300 \cdot 2^{32}) + 17612864$ , and so  $r_1 = 76300$  and  $r_0 = 17612864$ . The 14 most significant bits of  $r_0$  yield the value  $h_a(k) = 67$ .

Even though the multiply-shift method is fast, it doesn't provide any guarantee of good average-case performance. The universal hashing approach presented in the next section provides such a guarantee. A simple randomized variant of the multiply-shift method works well on the average, when the program begins by picking  $a$  as a randomly chosen odd integer.

### 11.3.2 Random hashing

Suppose that a malicious adversary chooses the keys to be hashed by some fixed hash function. Then the adversary can choose  $n$  keys that all hash to the same slot, yielding an average retrieval time of  $\Theta(n)$ . Any static hash function is vulnerable to such terrible worst-case behavior. The only effective way to improve the situation is to choose the hash function *randomly* in a way that is *independent* of the keys that are actually going to be stored. This approach is called **random hashing**. A special case of this approach, called **universal hashing**, can yield provably good performance on average when collisions are handled by chaining, no matter which keys the adversary chooses.

To use random hashing, at the beginning of program execution you select the hash function at random from a suitable family of functions. As in the case of quicksort, randomization guarantees that no single input always evokes worst-case behavior. Because you randomly select the hash function, the algorithm can behave differently on each execution, even for the same set of keys to be hashed, guaranteeing good average-case performance.

Let  $\mathcal{H}$  be a finite family of hash functions that map a given universe  $U$  of keys into the range  $\{0, 1, \dots, m-1\}$ . Such a family is said to be **universal** if for each pair of distinct keys  $k_1, k_2 \in U$ , the number of hash functions  $h \in \mathcal{H}$  for which  $h(k_1) = h(k_2)$  is at most  $|\mathcal{H}|/m$ . In other words, with a hash function randomly chosen from  $\mathcal{H}$ , the chance of a collision between distinct keys  $k_1$  and  $k_2$  is no more than the chance  $1/m$  of a collision if  $h(k_1)$  and  $h(k_2)$  were randomly and independently chosen from the set  $\{0, 1, \dots, m-1\}$ .

Independent uniform hashing is the same as picking a hash function uniformly at random from a family of  $m^n$  hash functions, each member of that family mapping the  $n$  keys to the  $m$  hash values in a different way.

Every independent uniform random family of hash function is universal, but the converse need not be true: consider the case where  $U = \{0, 1, \dots, m-1\}$  and the only hash function in the family is the identity function. The probability that two distinct keys collide is zero, even though each key is hashed to a fixed value.

The following corollary to Theorem 11.2 on page 279 says that universal hashing provides the desired payoff: it becomes impossible for an adversary to pick a sequence of operations that forces the worst-case running time.

**Corollary 11.3**

Using universal hashing and collision resolution by chaining in an initially empty table with  $m$  slots, it takes  $\Theta(s)$  expected time to handle any sequence of  $s$  INSERT, SEARCH, and DELETE operations containing  $n = O(m)$  INSERT operations.

**Proof** The INSERT and DELETE operations take constant time. Since the number  $n$  of insertions is  $O(m)$ , we have that  $\alpha = O(1)$ . Furthermore, the expected time for each SEARCH operation is  $O(1)$ , which can be seen by examining the proof of Theorem 11.2. That analysis depends only on collision probabilities, which are  $1/m$  for any pair  $k_1, k_2$  of keys by the choice of an independent uniform hash function in that theorem. Using a universal family of hash functions here instead of using independent uniform hashing changes the probability of collision from  $1/m$  to at most  $1/m$ . By linearity of expectation, therefore, the expected time for the entire sequence of  $s$  operations is  $O(s)$ . Since each operation takes  $\Omega(1)$  time, the  $\Theta(s)$  bound follows. ■

### 11.3.3 Achievable properties of random hashing

There is a rich literature on the properties a family  $\mathcal{H}$  of hash functions can have, and how they relate to the efficiency of hashing. We summarize a few of the most interesting ones here.

Let  $\mathcal{H}$  be a family of hash functions, each with domain  $U$  and range  $\{0, 1, \dots, m-1\}$ , and let  $h$  be any hash function that is picked uniformly at random from  $\mathcal{H}$ . The probabilities mentioned are probabilities over the picks of  $h$ .

- The family  $\mathcal{H}$  is **uniform** if for any key  $k$  in  $U$  and any slot  $q$  in the range  $\{0, 1, \dots, m-1\}$ , the probability that  $h(k) = q$  is  $1/m$ .
- The family  $\mathcal{H}$  is **universal** if for any distinct keys  $k_1$  and  $k_2$  in  $U$ , the probability that  $h(k_1) = h(k_2)$  is at most  $1/m$ .
- The family  $\mathcal{H}$  of hash functions is  **$\epsilon$ -universal** if for any distinct keys  $k_1$  and  $k_2$  in  $U$ , the probability that  $h(k_1) = h(k_2)$  is at most  $\epsilon$ . Therefore, a universal family of hash functions is also  $1/m$ -universal.<sup>2</sup>

---

<sup>2</sup> In the literature, a  $(c/m)$ -universal hash function is sometimes called  $c$ -universal or  $c$ -approximately universal. We'll stick with the notation  $(c/m)$ -universal.

- The family  $\mathcal{H}$  is ***d-independent*** if for any distinct keys  $k_1, k_2, \dots, k_d$  in  $U$  and any slots  $q_1, q_2, \dots, q_d$ , not necessarily distinct, in  $\{0, 1, \dots, m-1\}$  the probability that  $h(k_i) = q_i$  for  $i = 1, 2, \dots, d$  is  $1/m^d$ .

Universal hash-function families are of particular interest, as they are the simplest type supporting provably efficient hash-table operations for any input data set. Many other interesting and desirable properties, such as those noted above, are also possible and allow for efficient specialized hash-table operations.

### 11.3.4 Designing a universal family of hash functions

This section presents two ways to design a universal (or  $\epsilon$ -universal) family of hash functions: one based on number theory and another based on a randomized variant of the multiply-shift method presented in Section 11.3.1. The first method is a bit easier to prove universal, but the second method is newer and faster in practice.

#### A universal family of hash functions based on number theory

We can design a universal family of hash functions using a little number theory. You may wish to refer to Chapter 31 if you are unfamiliar with basic concepts in number theory.

Begin by choosing a prime number  $p$  large enough so that every possible key  $k$  lies in the range 0 to  $p-1$ , inclusive. We assume here that  $p$  has a “reasonable” length. (See Section 11.3.5 for a discussion of methods for handling long input keys, such as variable-length strings.) Let  $\mathbb{Z}_p$  denote the set  $\{0, 1, \dots, p-1\}$ , and let  $\mathbb{Z}_p^*$  denote the set  $\{1, 2, \dots, p-1\}$ . Since  $p$  is prime, we can solve equations modulo  $p$  with the methods given in Chapter 31. Because the size of the universe of keys is greater than the number of slots in the hash table (otherwise, just use direct addressing), we have  $p > m$ .

Given any  $a \in \mathbb{Z}_p^*$  and any  $b \in \mathbb{Z}_p$ , define the hash function  $h_{ab}$  as a linear transformation followed by reductions modulo  $p$  and then modulo  $m$ :

$$h_{ab}(k) = ((ak + b) \bmod p) \bmod m. \quad (11.3)$$

For example, with  $p = 17$  and  $m = 6$ , we have

$$\begin{aligned} h_{3,4}(8) &= ((3 \cdot 8 + 4) \bmod 17) \bmod 6 \\ &= (28 \bmod 17) \bmod 6 \\ &= 11 \bmod 6 \\ &= 5. \end{aligned}$$

Given  $p$  and  $m$ , the family of all such hash functions is

$$\mathcal{H}_{pm} = \{h_{ab} : a \in \mathbb{Z}_p^* \text{ and } b \in \mathbb{Z}_p\}. \quad (11.4)$$

Each hash function  $h_{ab}$  maps  $\mathbb{Z}_p$  to  $\mathbb{Z}_m$ . This family of hash functions has the nice property that the size  $m$  of the output range (which is the size of the hash table) is arbitrary—it need not be prime. Since you can choose from among  $p - 1$  values for  $a$  and  $p$  values for  $b$ , the family  $\mathcal{H}_{pm}$  contains  $p(p - 1)$  hash functions.

**Theorem 11.4**

The family  $\mathcal{H}_{pm}$  of hash functions defined by equations (11.3) and (11.4) is universal.

**Proof** Consider two distinct keys  $k_1$  and  $k_2$  from  $\mathbb{Z}_p$ , so that  $k_1 \neq k_2$ . For a given hash function  $h_{ab}$ , let

$$\begin{aligned} r_1 &= (ak_1 + b) \bmod p, \\ r_2 &= (ak_2 + b) \bmod p. \end{aligned}$$

We first note that  $r_1 \neq r_2$ . Why? Since we have  $r_1 - r_2 = a(k_1 - k_2) \pmod{p}$ , it follows that  $r_1 \neq r_2$  because  $p$  is prime and both  $a$  and  $(k_1 - k_2)$  are nonzero modulo  $p$ . By Theorem 31.6 on page 908, their product must also be nonzero modulo  $p$ . Therefore, when computing any  $h_{ab} \in \mathcal{H}_{pm}$ , distinct inputs  $k_1$  and  $k_2$  map to distinct values  $r_1$  and  $r_2$  modulo  $p$ , and there are no collisions yet at the “mod  $p$  level.” Moreover, each of the possible  $p(p - 1)$  choices for the pair  $(a, b)$  with  $a \neq 0$  yields a *different* resulting pair  $(r_1, r_2)$  with  $r_1 \neq r_2$ , since we can solve for  $a$  and  $b$  given  $r_1$  and  $r_2$ :

$$\begin{aligned} a &= ((r_1 - r_2)((k_1 - k_2)^{-1} \bmod p)) \bmod p, \\ b &= (r_1 - ak_1) \bmod p, \end{aligned}$$

where  $((k_1 - k_2)^{-1} \bmod p)$  denotes the unique multiplicative inverse, modulo  $p$ , of  $k_1 - k_2$ . For each of the  $p$  possible values of  $r_1$ , there are only  $p - 1$  possible values of  $r_2$  that do not equal  $r_1$ , making only  $p(p - 1)$  possible pairs  $(r_1, r_2)$  with  $r_1 \neq r_2$ . Therefore, there is a one-to-one correspondence between pairs  $(a, b)$  with  $a \neq 0$  and pairs  $(r_1, r_2)$  with  $r_1 \neq r_2$ . Thus, for any given pair of distinct inputs  $k_1$  and  $k_2$ , if we pick  $(a, b)$  uniformly at random from  $\mathbb{Z}_p^* \times \mathbb{Z}_p$ , the resulting pair  $(r_1, r_2)$  is equally likely to be any pair of distinct values modulo  $p$ .

Therefore, the probability that distinct keys  $k_1$  and  $k_2$  collide is equal to the probability that  $r_1 = r_2 \pmod{m}$  when  $r_1$  and  $r_2$  are randomly chosen as distinct values modulo  $p$ . For a given value of  $r_1$ , of the  $p - 1$  possible remaining values for  $r_2$ , the number of values  $r_2$  such that  $r_2 \neq r_1$  and  $r_2 = r_1 \pmod{m}$  is at most

$$\begin{aligned} \left\lceil \frac{p}{m} \right\rceil - 1 &\leq \frac{p + m - 1}{m} - 1 \quad (\text{by inequality (3.7) on page 64}) \\ &= \frac{p - 1}{m}. \end{aligned}$$

The probability that  $r_2$  collides with  $r_1$  when reduced modulo  $m$  is at most  $((p-1)/m)/(p-1) = 1/m$ , since  $r_2$  is equally likely to be any of the  $p-1$  values in  $\mathbb{Z}_p$  that are different from  $r_1$ , but at most  $(p-1)/m$  of those values are equivalent to  $r_1$  modulo  $m$ .

Therefore, for any pair of distinct values  $k_1, k_2 \in \mathbb{Z}_p$ ,

$$\Pr \{h_{ab}(k_1) = h_{ab}(k_2)\} \leq 1/m ,$$

so that  $\mathcal{H}_{pm}$  is indeed universal. ■

### A $2/m$ -universal family of hash functions based on the multiply-shift method

We recommend that in practice you use the following hash-function family based on the multiply-shift method. It is exceptionally efficient and (although we omit the proof) provably  $2/m$ -universal. Define  $\mathcal{H}$  to be the family of multiply-shift hash functions with odd constants  $a$ :

$$\mathcal{H} = \{h_a : a \text{ is odd}, 1 \leq a < m, \text{ and } h_a \text{ is defined by equation (11.2)}\} . \quad (11.5)$$

#### **Theorem 11.5**

The family of hash functions  $\mathcal{H}$  given by equation (11.5) is  $2/m$ -universal. ■

That is, the probability that any two distinct keys collide is at most  $2/m$ . In many practical situations, the speed of computing the hash function more than compensates for the higher upper bound on the probability that two distinct keys collide when compared with a universal hash function.

### 11.3.5 Hashing long inputs such as vectors or strings

Sometimes hash function inputs are so long that they cannot be easily encoded modulo a reasonably sized prime number  $p$  or encoded within a single word of, say, 64 bits. As an example, consider the class of vectors, such as vectors of 8-bit bytes (which is how strings in many programming languages are stored). A vector might have an arbitrary nonnegative length, in which case the length of the input to the hash function may vary from input to input.

#### **Number-theoretic approaches**

One way to design good hash functions for variable-length inputs is to extend the ideas used in Section 11.3.4 to design universal hash functions. Exercise 11.3-6 explores one such approach.

### Cryptographic hashing

Another way to design a good hash function for variable-length inputs is to use a hash function designed for cryptographic applications. *Cryptographic hash functions* are complex pseudorandom functions, designed for applications requiring properties beyond those needed here, but are robust, widely implemented, and usable as hash functions for hash tables.

A cryptographic hash function takes as input an arbitrary byte string and returns a fixed-length output. For example, the NIST standard deterministic cryptographic hash function SHA-256 [346] produces a 256-bit (32-byte) output for any input.

Some chip manufacturers include instructions in their CPU architectures to provide fast implementations of some cryptographic functions. Of particular interest are instructions that efficiently implement rounds of the Advanced Encryption Standard (AES), the “AES-NI” instructions. These instructions execute in a few tens of nanoseconds, which is generally fast enough for use with hash tables. A message authentication code such as CBC-MAC based on AES and the use of the AES-NI instructions could be a useful and efficient hash function. We don’t pursue the potential use of specialized instruction sets further here.

Cryptographic hash functions are useful because they provide a way of implementing an approximate version of a random oracle. As noted earlier, a random oracle is equivalent to an independent uniform hash function family. From a theoretical point of view, a random oracle is an unachievable ideal: a deterministic function that provides a randomly selected output for each input. Because it is deterministic, it provides the same output if queried again for the same input. From a practical point of view, constructions of hash function families based on cryptographic hash functions are sensible substitutes for random oracles.

There are many ways to use a cryptographic hash function as a hash function. For example, we could define

$$h(k) = \text{SHA-256}(k) \bmod m .$$

To define a family of such hash functions one may prepend a “salt” string  $a$  to the input before hashing it, as in

$$h_a(k) = \text{SHA-256}(a \parallel k) \bmod m ,$$

where  $a \parallel k$  denotes the string formed by concatenating the strings  $a$  and  $k$ . The literature on message authentication codes (MACs) provides additional approaches.

Cryptographic approaches to hash-function design are becoming more practical as computers arrange their memories in hierarchies of differing capacities and speeds. Section 11.5 discusses one hash-function design based on the RC6 encryption method.



**Exercises****11.3-1**

You wish to search a linked list of length  $n$ , where each element contains a key  $k$  along with a hash value  $h(k)$ . Each key is a long character string. How might you take advantage of the hash values when searching the list for an element with a given key?

**11.3-2**

You hash a string of  $r$  characters into  $m$  slots by treating it as a radix-128 number and then using the division method. You can represent the number  $m$  as a 32-bit computer word, but the string of  $r$  characters, treated as a radix-128 number, takes many words. How can you apply the division method to compute the hash value of the character string without using more than a constant number of words of storage outside the string itself?

**11.3-3**

Consider a version of the division method in which  $h(k) = k \bmod m$ , where  $m = 2^p - 1$  and  $k$  is a character string interpreted in radix  $2^p$ . Show that if string  $x$  can be converted to string  $y$  by permuting its characters, then  $x$  and  $y$  hash to the same value. Give an example of an application in which this property would be undesirable in a hash function.

**11.3-4**

Consider a hash table of size  $m = 1000$  and a corresponding hash function  $h(k) = \lfloor m(kA \bmod 1) \rfloor$  for  $A = (\sqrt{5} - 1)/2$ . Compute the locations to which the keys 61, 62, 63, 64, and 65 are mapped.

**★ 11.3-5**

Show that any  $\epsilon$ -universal family  $\mathcal{H}$  of hash functions from a finite set  $U$  to a finite set  $Q$  has  $\epsilon \geq 1/|Q| - 1/|U|$ .

**★ 11.3-6**

Let  $U$  be the set of  $d$ -tuples of values drawn from  $\mathbb{Z}_p$ , and let  $Q = \mathbb{Z}_p$ , where  $p$  is prime. Define the hash function  $h_b : U \rightarrow Q$  for  $b \in \mathbb{Z}_p$  on an input  $d$ -tuple  $\langle a_0, a_1, \dots, a_{d-1} \rangle$  from  $U$  as

$$h_b(\langle a_0, a_1, \dots, a_{d-1} \rangle) = \left( \sum_{j=0}^{d-1} a_j b^j \right) \bmod p,$$

and let  $\mathcal{H} = \{h_b : b \in \mathbb{Z}_p\}$ . Argue that  $\mathcal{H}$  is  $\epsilon$ -universal for  $\epsilon = (d-1)/p$ . (Hint: See Exercise 31.4-4.)



---

## 11.4 Open addressing

This section describes open addressing, a method for collision resolution that, unlike chaining, does not make use of storage outside of the hash table itself. In *open addressing*, all elements occupy the hash table itself. That is, each table entry contains either an element of the dynamic set or NIL. No lists or elements are stored outside the table, unlike in chaining. Thus, in open addressing, the hash table can “fill up” so that no further insertions can be made. One consequence is that the load factor  $\alpha$  can never exceed 1.

Collisions are handled as follows: when a new element is to be inserted into the table, it is placed in its “first-choice” location if possible. If that location is already occupied, the new element is placed in its “second-choice” location. The process continues until an empty slot is found in which to place the new element. Different elements have different preference orders for the locations.

To search for an element, systematically examine the preferred table slots for that element, in order of decreasing preference, until either you find the desired element or you find an empty slot and thus verify that the element is not in the table.

Of course, you could use chaining and store the linked lists inside the hash table, in the otherwise unused hash-table slots (see Exercise 11.2-4), but the advantage of open addressing is that it avoids pointers altogether. Instead of following pointers, you compute the sequence of slots to be examined. The memory freed by not storing pointers provides the hash table with a larger number of slots in the same amount of memory, potentially yielding fewer collisions and faster retrieval.

To perform insertion using open addressing, successively examine, or *probe*, the hash table until you find an empty slot in which to put the key. Instead of being fixed in the order  $0, 1, \dots, m-1$  (which implies a  $\Theta(n)$  search time), the sequence of positions probed depends upon the key being inserted. To determine which slots to probe, the hash function includes the probe number (starting from 0) as a second input. Thus, the hash function becomes

$$h : U \times \{0, 1, \dots, m-1\} \rightarrow \{0, 1, \dots, m-1\} .$$

Open addressing requires that for every key  $k$ , the *probe sequence*  $\langle h(k, 0), h(k, 1), \dots, h(k, m-1) \rangle$  be a permutation of  $\{0, 1, \dots, m-1\}$ , so that every hash-table position is eventually considered as a slot for a new key as the table fills up. The HASH-INSERT procedure on the following page assumes that the elements in the hash table  $T$  are keys with no satellite information: the key  $k$  is identical to the element containing key  $k$ . Each slot contains either a key or NIL (if the slot is empty). The HASH-INSERT procedure takes as input a hash table  $T$  and a key  $k$

that is assumed to be not already present in the hash table. It either returns the slot number where it stores key  $k$  or flags an error because the hash table is already full.

```

HASH-INSERT( $T, k$ )
1   $i = 0$ 
2  repeat
3       $q = h(k, i)$ 
4      if  $T[q] == \text{NIL}$ 
5           $T[q] = k$ 
6          return  $q$ 
7      else  $i = i + 1$ 
8  until  $i == m$ 
9  error “hash table overflow”

HASH-SEARCH( $T, k$ )
1   $i = 0$ 
2  repeat
3       $q = h(k, i)$ 
4      if  $T[q] == k$ 
5          return  $q$ 
6       $i = i + 1$ 
7  until  $T[q] == \text{NIL}$  or  $i == m$ 
8  return NIL

```

The algorithm for searching for key  $k$  probes the same sequence of slots that the insertion algorithm examined when key  $k$  was inserted. Therefore, the search can terminate (unsuccessfully) when it finds an empty slot, since  $k$  would have been inserted there and not later in its probe sequence. The procedure HASH-SEARCH takes as input a hash table  $T$  and a key  $k$ , returning  $q$  if it finds that slot  $q$  contains key  $k$ , or NIL if key  $k$  is not present in table  $T$ .

Deletion from an open-address hash table is tricky. When you delete a key from slot  $q$ , it would be a mistake to mark that slot as empty by simply storing NIL in it. If you did, you might be unable to retrieve any key  $k$  for which slot  $q$  was probed and found occupied when  $k$  was inserted. One way to solve this problem is by marking the slot, storing in it the special value DELETED instead of NIL. The HASH-INSERT procedure then has to treat such a slot as empty so that it can insert a new key there. The HASH-SEARCH procedure passes over DELETED values while searching, since slots containing DELETED were filled when the key being searched for was inserted. Using the special value DELETED, however, means that search times no longer depend on the load factor  $\alpha$ , and for this reason chaining is

frequently selected as a collision resolution technique when keys must be deleted. There is a simple special case of open addressing, linear probing, that avoids the need to mark slots with DELETED. Section 11.5.1 shows how to delete from a hash table when using linear probing.

In our analysis, we assume *independent uniform permutation hashing* (also confusingly known as *uniform hashing* in the literature): the probe sequence of each key is equally likely to be any of the  $m!$  permutations of  $\langle 0, 1, \dots, m-1 \rangle$ . Independent uniform permutation hashing generalizes the notion of independent uniform hashing defined earlier to a hash function that produces not just a single slot number, but a whole probe sequence. True independent uniform permutation hashing is difficult to implement, however, and in practice suitable approximations (such as double hashing, defined below) are used.

We'll examine both double hashing and its special case, linear probing. These techniques guarantee that  $\langle h(k, 0), h(k, 1), \dots, h(k, m-1) \rangle$  is a permutation of  $\langle 0, 1, \dots, m-1 \rangle$  for each key  $k$ . (Recall that the second parameter to the hash function  $h$  is the probe number.) Neither double hashing nor linear probing meets the assumption of independent uniform permutation hashing, however. Double hashing cannot generate more than  $m^2$  different probe sequences (instead of the  $m!$  that independent uniform permutation hashing requires). Nonetheless, double hashing has a large number of possible probe sequences and, as you might expect, seems to give good results. Linear probing is even more restricted, capable of generating only  $m$  different probe sequences.


### Double hashing

Double hashing offers one of the best methods available for open addressing because the permutations produced have many of the characteristics of randomly chosen permutations. *Double hashing* uses a hash function of the form

$$h(k, i) = (h_1(k) + i h_2(k)) \bmod m ,$$

where both  $h_1$  and  $h_2$  are *auxiliary hash functions*. The initial probe goes to position  $T[h_1(k)]$ , and successive probe positions are offset from previous positions by the amount  $h_2(k)$ , modulo  $m$ . Thus, the probe sequence here depends in two ways upon the key  $k$ , since the initial probe position  $h_1(k)$ , the step size  $h_2(k)$ , or both, may vary. Figure 11.5 gives an example of insertion by double hashing.

In order for the entire hash table to be searched, the value  $h_2(k)$  must be relatively prime to the hash-table size  $m$ . (See Exercise 11.4-5.) A convenient way to ensure this condition is to let  $m$  be an exact power of 2 and to design  $h_2$  so that it always produces an odd number. Another way is to let  $m$  be prime and to design  $h_2$  so that it always returns a positive integer less than  $m$ . For example, you



**Figure 11.5** Insertion by double hashing. The hash table has size 13 with  $h_1(k) = k \bmod 13$  and  $h_2(k) = 1 + (k \bmod 11)$ . Since  $14 = 1 \pmod{13}$  and  $14 = 3 \pmod{11}$ , the key 14 goes into empty slot 9, after slots 1 and 5 are examined and found to be occupied.

could choose  $m$  prime and let

$$h_1(k) = k \bmod m ,$$

$$h_2(k) = 1 + (k \bmod m') ,$$

where  $m'$  is chosen to be slightly less than  $m$  (say,  $m - 1$ ). For example, if  $k = 123456$ ,  $m = 701$ , and  $m' = 700$ , then  $h_1(k) = 80$  and  $h_2(k) = 257$ , so that the first probe goes to position 80, and successive probes examine every 257th slot (modulo  $m$ ) until the key has been found or every slot has been examined.

Although values of  $m$  other than primes or exact powers of 2 can in principle be used with double hashing, in practice it becomes more difficult to efficiently generate  $h_2(k)$  (other than choosing  $h_2(k) = 1$ , which gives linear probing) in a way that ensures that it is relatively prime to  $m$ , in part because the relative density  $\phi(m)/m$  of such numbers for general  $m$  may be small (see equation (31.25) on page 921).

When  $m$  is prime or an exact power of 2, double hashing produces  $\Theta(m^2)$  probe sequences, since each possible  $(h_1(k), h_2(k))$  pair yields a distinct probe sequence. As a result, for such values of  $m$ , double hashing appears to perform close to the “ideal” scheme of independent uniform permutation hashing.

### Linear probing

**Linear probing**, a special case of double hashing, is the simplest open-addressing approach to resolving collisions. As with double hashing, an auxiliary hash function  $h_1$  determines the first probe position  $h_1(k)$  for inserting an element. If slot  $T[h_1(k)]$  is already occupied, probe the next position  $T[h_1(k) + 1]$ . Keep going as necessary, on up to slot  $T[m - 1]$ , and then wrap around to slots  $T[0]$ ,  $T[1]$ , and so on, but never going past slot  $T[h_1(k) - 1]$ . To view linear probing as a special case of double hashing, just set the double-hashing step function  $h_2$  to be fixed at 1:  $h_2(k) = 1$  for all  $k$ . That is, the hash function is

$$h(k, i) = (h_1(k) + i) \bmod m \quad (11.6)$$

for  $i = 0, 1, \dots, m - 1$ . The value of  $h_1(k)$  determines the entire probe sequence, and so assuming that  $h_1(k)$  can take on any value in  $\{0, 1, \dots, m - 1\}$ , linear probing allows only  $m$  distinct probe sequences.

We'll revisit linear probing in Section 11.5.1.

### Analysis of open-address hashing

As in our analysis of chaining in Section 11.2, we analyze open addressing in terms of the load factor  $\alpha = n/m$  of the hash table. With open addressing, at most one element occupies each slot, and thus  $n \leq m$ , which implies  $\alpha \leq 1$ . The analysis below requires  $\alpha$  to be strictly less than 1, and so we assume that at least one slot is empty. Because deleting from an open-address hash table does not really free up a slot, we assume as well that no deletions occur.

For the hash function, we assume independent uniform permutation hashing. In this idealized scheme, the probe sequence  $\langle h(k, 0), h(k, 1), \dots, h(k, m - 1) \rangle$  used to insert or search for each key  $k$  is equally likely to be any permutation of  $\langle 0, 1, \dots, m - 1 \rangle$ . Of course, any given key has a unique fixed probe sequence associated with it. What we mean here is that, considering the probability distribution on the space of keys and the operation of the hash function on the keys, each possible probe sequence is equally likely.

We now analyze the expected number of probes for hashing with open addressing under the assumption of independent uniform permutation hashing, beginning with the expected number of probes made in an unsuccessful search (assuming, as stated above, that  $\alpha < 1$ ).

The bound proven, of  $1/(1 - \alpha) = 1 + \alpha + \alpha^2 + \alpha^3 + \dots$ , has an intuitive interpretation. The first probe always occurs. With probability approximately  $\alpha$ , the first probe finds an occupied slot, so that a second probe happens. With probability approximately  $\alpha^2$ , the first two slots are occupied so that a third probe ensues, and so on.

**Theorem 11.6**

Given an open-address hash table with load factor  $\alpha = n/m < 1$ , the expected number of probes in an unsuccessful search is at most  $1/(1 - \alpha)$ , assuming independent uniform permutation hashing and no deletions.

**Proof** In an unsuccessful search, every probe but the last accesses an occupied slot that does not contain the desired key, and the last slot probed is empty. Let the random variable  $X$  denote the number of probes made in an unsuccessful search, and define the event  $A_i$ , for  $i = 1, 2, \dots$ , as the event that an  $i$ th probe occurs and it is to an occupied slot. Then the event  $\{X \geq i\}$  is the intersection of events  $A_1 \cap A_2 \cap \dots \cap A_{i-1}$ . We bound  $\Pr\{X \geq i\}$  by bounding  $\Pr\{A_1 \cap A_2 \cap \dots \cap A_{i-1}\}$ . By Exercise C.2-5 on page 1190,

$$\begin{aligned} \Pr\{A_1 \cap A_2 \cap \dots \cap A_{i-1}\} &= \Pr\{A_1\} \cdot \Pr\{A_2 \mid A_1\} \cdot \Pr\{A_3 \mid A_1 \cap A_2\} \cdots \\ &\quad \Pr\{A_{i-1} \mid A_1 \cap A_2 \cap \dots \cap A_{i-2}\}. \end{aligned}$$

Since there are  $n$  elements and  $m$  slots,  $\Pr\{A_1\} = n/m$ . For  $j > 1$ , the probability that there is a  $j$ th probe and it is to an occupied slot, given that the first  $j - 1$  probes were to occupied slots, is  $(n - j + 1)/(m - j + 1)$ . This probability follows because the  $j$ th probe would be finding one of the remaining  $(n - (j - 1))$  elements in one of the  $(m - (j - 1))$  unexamined slots, and by the assumption of independent uniform permutation hashing, the probability is the ratio of these quantities. Since  $n < m$  implies that  $(n - j)/(m - j) \leq n/m$  for all  $j$  in the range  $0 \leq j < m$ , it follows that for all  $i$  in the range  $1 \leq i \leq m$ , we have

$$\begin{aligned} \Pr\{X \geq i\} &= \frac{n}{m} \cdot \frac{n-1}{m-1} \cdot \frac{n-2}{m-2} \cdots \frac{n-i+2}{m-i+2} \\ &\leq \left(\frac{n}{m}\right)^{i-1} \\ &= \alpha^{i-1}. \end{aligned}$$

The product in the first line has  $i - 1$  factors. When  $i = 1$ , the product is 1, the identity for multiplication, and we get  $\Pr\{X \geq 1\} = 1$ , which makes sense, since there must always be at least 1 probe. If each of the first  $n$  probes is to an occupied slot, then all occupied slots have been probed. Then, the  $(n + 1)$ st probe must be to an empty slot, which gives  $\Pr\{X \geq i\} = 0$  for  $i > n + 1$ . Now, we use equation (C.28) on page 1193 to bound the expected number of probes:

$$\begin{aligned} E[X] &= \sum_{i=1}^{\infty} \Pr\{X \geq i\} \\ &= \sum_{i=1}^{n+1} \Pr\{X \geq i\} + \sum_{i>n+1} \Pr\{X \geq i\} \end{aligned}$$

$$\begin{aligned}
&\leq \sum_{i=1}^{\infty} \alpha^{i-1} + 0 \\
&= \sum_{i=0}^{\infty} \alpha^i \\
&= \frac{1}{1-\alpha} \quad (\text{by equation (A.7) on page 1142 because } 0 \leq \alpha < 1) . \quad \blacksquare
\end{aligned}$$

If  $\alpha$  is a constant, Theorem 11.6 predicts that an unsuccessful search runs in  $O(1)$  time. For example, if the hash table is half full, the average number of probes in an unsuccessful search is at most  $1/(1 - .5) = 2$ . If it is 90% full, the average number of probes is at most  $1/(1 - .9) = 10$ .

Theorem 11.6 yields almost immediately how well the HASH-INSERT procedure performs.

### **Corollary 11.7**

Inserting an element into an open-address hash table with load factor  $\alpha$ , where  $\alpha < 1$ , requires at most  $1/(1 - \alpha)$  probes on average, assuming independent uniform permutation hashing and no deletions.

**Proof** An element is inserted only if there is room in the table, and thus  $\alpha < 1$ . Inserting a key requires an unsuccessful search followed by placing the key into the first empty slot found. Thus, the expected number of probes is at most  $1/(1 - \alpha)$ . ■

It takes a little more work to compute the expected number of probes for a successful search.

### **Theorem 11.8**

Given an open-address hash table with load factor  $\alpha < 1$ , the expected number of probes in a successful search is at most

$$\frac{1}{\alpha} \ln \frac{1}{1-\alpha} ,$$

assuming independent uniform permutation hashing with no deletions and assuming that each key in the table is equally likely to be searched for.

**Proof** A search for a key  $k$  reproduces the same probe sequence as when the element with key  $k$  was inserted. If  $k$  was the  $(i + 1)$ st key inserted into the hash table, then the load factor at the time it was inserted was  $i/m$ , and so by Corollary 11.7, the expected number of probes made in a search for  $k$  is at most  $1/(1 - i/m) = m/(m - i)$ . Averaging over all  $n$  keys in the hash table gives us

the expected number of probes in a successful search:

$$\begin{aligned}
 \frac{1}{n} \sum_{i=0}^{n-1} \frac{m}{m-i} &= \frac{m}{n} \sum_{i=0}^{n-1} \frac{1}{m-i} \\
 &= \frac{1}{\alpha} \sum_{k=m-n+1}^m \frac{1}{k} \\
 &\leq \frac{1}{\alpha} \int_{m-n}^m \frac{1}{x} dx \quad (\text{by inequality (A.19) on page 1150}) \\
 &= \frac{1}{\alpha} (\ln m - \ln(m-n)) \\
 &= \frac{1}{\alpha} \ln \frac{m}{m-n} \\
 &= \frac{1}{\alpha} \ln \frac{1}{1-\alpha}. \quad \blacksquare
 \end{aligned}$$

If the hash table is half full, the expected number of probes in a successful search is less than 1.387. If the hash table is 90% full, the expected number of probes is less than 2.559. If  $\alpha = 1$ , then in an unsuccessful search, all  $m$  slots must be probed. Exercise 11.4-4 asks you to analyze a successful search when  $\alpha = 1$ .

## Exercises

### 11.4-1

Consider inserting the keys 10, 22, 31, 4, 15, 28, 17, 88, 59 into a hash table of length  $m = 11$  using open addressing. Illustrate the result of inserting these keys using linear probing with  $h(k, i) = (k + i) \bmod m$  and using double hashing with  $h_1(k) = k$  and  $h_2(k) = 1 + (k \bmod (m - 1))$ .

### 11.4-2

Write pseudocode for HASH-DELETE that fills the deleted key's slot with the special value DELETED, and modify HASH-SEARCH and HASH-INSERT as needed to handle DELETED.

### 11.4-3

Consider an open-address hash table with independent uniform permutation hashing and no deletions. Give upper bounds on the expected number of probes in an unsuccessful search and on the expected number of probes in a successful search when the load factor is  $3/4$  and when it is  $7/8$ .



**11.4-4**

Show that the expected number of probes required for a successful search when  $\alpha = 1$  (that is, when  $n = m$ ), is  $H_m$ , the  $m$ th harmonic number.

**★ 11.4-5**

Show that, with double hashing, if  $m$  and  $h_2(k)$  have greatest common divisor  $d \geq 1$  for some key  $k$ , then an unsuccessful search for key  $k$  examines  $(1/d)$ th of the hash table before returning to slot  $h_1(k)$ . Thus, when  $d = 1$ , so that  $m$  and  $h_2(k)$  are relatively prime, the search may examine the entire hash table. (*Hint:* See Chapter 31.)

**★ 11.4-6**

Consider an open-address hash table with a load factor  $\alpha$ . Approximate the nonzero value  $\alpha$  for which the expected number of probes in an unsuccessful search equals twice the expected number of probes in a successful search. Use the upper bounds given by Theorems 11.6 and 11.8 for these expected numbers of probes.

---

## 11.5 Practical considerations

Efficient hash table algorithms are not only of theoretical interest, but also of immense practical importance. Constant factors can matter. For this reason, this section discusses two aspects of modern CPUs that are not included in the standard RAM model presented in Section 2.2:

**Memory hierarchies:** The memory of modern CPUs has a number of levels, from the fast registers, through one or more levels of *cache memory*, to the main-memory level. Each successive level stores more data than the previous level, but access is slower. As a consequence, a complex computation (such as a complicated hash function) that works entirely within the fast registers can take less time than a single read operation from main memory. Furthermore, cache memory is organized in *cache blocks* of (say) 64 bytes each, which are always fetched together from main memory. There is a substantial benefit for ensuring that memory usage is local: reusing the same cache block is much more efficient than fetching a different cache block from main memory.

The standard RAM model measures efficiency of a hash-table operation by counting the number of hash-table slots probed. In practice, this metric is only a crude approximation to the truth, since once a cache block is in the cache, successive probes to that cache block are much faster than probes that must access main memory.

**Advanced instruction sets:** Modern CPUs may have sophisticated instruction sets that implement advanced primitives useful for encryption or other forms of cryptography. These instructions may be useful in the design of exceptionally efficient hash functions.

Section 11.5.1 discusses linear probing, which becomes the collision-resolution method of choice in the presence of a memory hierarchy. Section 11.5.2 suggests how to construct “advanced” hash functions based on cryptographic primitives, suitable for use on computers with hierarchical memory models.

### 11.5.1 Linear probing

Linear probing is often disparaged because of its poor performance in the standard RAM model. But linear probing excels for hierarchical memory models, because successive probes are usually to the same cache block of memory.

#### Deletion with linear probing

Another reason why linear probing is often not used in practice is that deletion seems complicated or impossible without using the special DELETED value. Yet we’ll now see that deletion from a hash table based on linear probing is not all that difficult, even without the DELETED marker. The deletion procedure works for linear probing, but not for open-address probing in general, because with linear probing keys all follow the same simple cyclic probing sequence (albeit with different starting points).

The deletion procedure relies on an “inverse” function to the linear-probing hash function  $h(k, i) = (h_1(k) + i) \bmod m$ , which maps a key  $k$  and a probe number  $i$  to a slot number in the hash table. The inverse function  $g$  maps a key  $k$  and a slot number  $q$ , where  $0 \leq q < m$ , to the probe number that reaches slot  $q$ :

$$g(k, q) = (q - h_1(k)) \bmod m.$$

If  $h(k, i) = q$ , then  $g(k, q) = i$ , and so  $h(k, g(k, q)) = q$ .

The procedure LINEAR-PROBING-HASH-DELETE on the facing page deletes the key stored in position  $q$  from hash table  $T$ . Figure 11.6 shows how it works. The procedure first deletes the key in position  $q$  by setting  $T[q]$  to NIL in line 2. It then searches for a slot  $q'$  (if any) that contains a key that should be moved to the slot  $q$  just vacated by key  $k$ . Line 9 asks the critical question: does the key  $k'$  in slot  $q'$  need to be moved to the vacated slot  $q$  in order to preserve the accessibility of  $k'$ ? If  $g(k', q) < g(k', q')$ , then during the insertion of  $k'$  into the table, slot  $q$  was examined but found to be already occupied. But now slot  $q$ , where a search will look for  $k'$ , is empty. In this case, key  $k'$  moves to slot  $q$  in line 10, and the

0		0	
1		1	
2	82	2	82
3	43	3	93
4	74	4	74
5	93	5	92
6	92	6	
7		7	
8	18	8	18
9	38	9	38
(a)		(b)	

**Figure 11.6** Deletion in a hash table that uses linear probing. The hash table has size 10 with  $h_1(k) = k \bmod 10$ . **(a)** The hash table after inserting keys in the order 74, 43, 93, 18, 82, 38, 92. **(b)** The hash table after deleting the key 43 from slot 3. Key 93 moves up to slot 3 to keep it accessible, and then key 92 moves up to slot 5 just vacated by key 93. No other keys need to be moved.

search continues, to see whether any later key also needs to be moved to the slot  $q'$  that was just freed up when  $k'$  moved.

#### LINEAR-PROBING-HASH-DELETE( $T, q$ )

```

1  while TRUE
2       $T[q] = \text{NIL}$                                 // make slot  $q$  empty
3       $q' = q$                                        // starting point for search
4      repeat
5           $q' = (q' + 1) \bmod m$                     // next slot number with linear probing
6           $k' = T[q']$                                // next key to try to move
7          if  $k' == \text{NIL}$ 
8              return                               // return when an empty slot is found
9      until  $g(k', q) < g(k', q')$                  // was empty slot  $q$  probed before  $q'$ ?
10      $T[q] = k'$                                     // move  $k'$  into slot  $q$ 
11      $q = q'$                                        // free up slot  $q'$ 
```

#### Analysis of linear probing

Linear probing is popular to implement, but it exhibits a phenomenon known as *primary clustering*. Long runs of occupied slots build up, increasing the average

search time. Clusters arise because an empty slot preceded by  $i$  full slots gets filled next with probability  $(i + 1)/m$ . Long runs of occupied slots tend to get longer, and the average search time increases.

In the standard RAM model, primary clustering is a problem, and general double hashing usually performs better than linear probing. By contrast, in a hierarchical memory model, primary clustering is a beneficial property, as elements are often stored together in the same cache block. Searching proceeds through one cache block before advancing to search the next cache block. With linear probing, the running time for a key  $k$  of HASH-INSERT, HASH-SEARCH, or LINEAR-PROBING-HASH-DELETE is at most proportional to the distance from  $h_1(k)$  to the next empty slot.

The following theorem is due to Pagh et al. [351]. A more recent proof is given by Thorup [438]. We omit the proof here. The need for 5-independence is by no means obvious; see the cited proofs.

**Theorem 11.9**

If  $h_1$  is 5-independent and  $\alpha \leq 2/3$ , then it takes expected constant time to search for, insert, or delete a key in a hash table using linear probing. ■

(Indeed, the expected operation time is  $O(1/\epsilon^2)$  for  $\alpha = 1 - \epsilon$ .)

★ **11.5.2 Hash functions for hierarchical memory models**

This section illustrates an approach for designing efficient hash tables in a modern computer system having a memory hierarchy.

Because of the memory hierarchy, linear probing is a good choice for resolving collisions, as probe sequences are sequential and tend to stay within cache blocks. But linear probing is most efficient when the hash function is complex (for example, 5-independent as in Theorem 11.9). Fortunately, having a memory hierarchy means that complex hash functions can be implemented efficiently.

As noted in Section 11.3.5, one approach is to use a cryptographic hash function such as SHA-256. Such functions are complex and sufficiently random for hash table applications. On machines with specialized instructions, cryptographic functions can be quite efficient.

Instead, we present here a simple hash function based only on addition, multiplication, and swapping the halves of a word. This function can be implemented entirely within the fast registers, and on a machine with a memory hierarchy, its latency is small compared with the time taken to access a random slot of the hash table. It is related to the RC6 encryption algorithm and can for practical purposes be considered a “random oracle.”

### The wee hash function

Let  $w$  denote the word size of the machine (e.g.,  $w = 64$ ), assumed to be even, and let  $a$  and  $b$  be  $w$ -bit unsigned (nonnegative) integers such that  $a$  is odd. Let  $\text{swap}(x)$  denote the  $w$ -bit result of swapping the two  $w/2$ -bit halves of  $w$ -bit input  $x$ . That is,

$$\text{swap}(x) = (x \ggg (w/2)) + (x \lll (w/2))$$

where “ $\ggg$ ” is “logical right shift” (as in equation (11.2)) and “ $\lll$ ” is “left shift.” Define

$$f_a(k) = \text{swap}((2k^2 + ak) \bmod 2^w).$$

Thus, to compute  $f_a(k)$ , evaluate the quadratic function  $2k^2 + ak$  modulo  $2^w$  and then swap the left and right halves of the result.

Let  $r$  denote a desired number of “rounds” for the computation of the hash function. We’ll use  $r = 4$ , but the hash function is well defined for any nonnegative  $r$ . Denote by  $f_a^{(r)}(k)$  the result of iterating  $f_a$  a total of  $r$  times (that is,  $r$  rounds) starting with input value  $k$ . For any odd  $a$  and any  $r \geq 0$ , the function  $f_a^{(r)}$ , although complicated, is one-to-one (see Exercise 11.5-1). A cryptographer would view  $f_a^{(r)}$  as a simple block cipher operating on  $w$ -bit input blocks, with  $r$  rounds and key  $a$ .

We first define the wee hash function  $h$  for short inputs, where by “short” we means “whose length  $t$  is at most  $w$ -bits,” so that the input fits within one computer word. We would like inputs of different lengths to be hashed differently. The *wee hash function*  $h_{a,b,t,r}(k)$  for parameters  $a$ ,  $b$ , and  $r$  on  $t$ -bit input  $k$  is defined as

$$h_{a,b,t,r}(k) = (f_{a+2t}^{(r)}(k + b)) \bmod m. \quad (11.7)$$

That is, the hash value for  $t$ -bit input  $k$  is obtained by applying  $f_{a+2t}^{(r)}$  to  $k + b$ , then taking the final result modulo  $m$ . Adding the value  $b$  provides hash-dependent randomization of the input, in a way that ensures that for variable-length inputs the 0-length input does not have a fixed hash value. Adding the value  $2t$  to  $a$  ensures that the hash function acts differently for inputs of different lengths. (We use  $2t$  rather than  $t$  to ensure that the key  $a + 2t$  is odd if  $a$  is odd.) We call this hash function “wee” because it uses a tiny amount of memory—more precisely, it can be implemented efficiently using only the computer’s fast registers. (This hash function does not have a name in the literature; it is a variant we developed for this textbook.)

### Speed of the wee hash function

It is surprising how much efficiency can be bought with locality. Experiments (unpublished, by the authors) suggest that evaluating the wee hash function takes less

time than probing a *single* randomly chosen slot in a hash table. These experiments were run on a laptop (2019 MacBook Pro) with  $w = 64$  and  $a = 123$ . For large hash tables, evaluating the wee hash function was 2 to 10 times faster than performing a single probe of the hash table.

### The wee hash function for variable-length inputs

Sometimes inputs are long—more than one  $w$ -bit word in length—or have variable length, as discussed in Section 11.3.5. We can extend the wee hash function, defined above for inputs that are at most single  $w$ -bit word in length, to handle long or variable-length inputs. Here is one method for doing so.

Suppose that an input  $k$  has length  $t$  (measured in bits). Break  $k$  into a sequence  $\langle k_1, k_2, \dots, k_u \rangle$  of  $w$ -bit words, where  $u = \lceil t/w \rceil$ ,  $k_1$  contains the least-significant  $w$  bits of  $k$ , and  $k_u$  contains the most significant bits. If  $t$  is not a multiple of  $w$ , then  $k_u$  contains fewer than  $w$  bits, in which case, pad out the unused high-order bits of  $k_u$  with 0-bits. Define the function *chop* to return a sequence of the  $w$ -bit words in  $k$ :

$$\text{chop}(k) = \langle k_1, k_2, \dots, k_u \rangle.$$

The most important property of the chop operation is that it is one-to-one, given  $t$ : for any two  $t$ -bit keys  $k$  and  $k'$ , if  $k \neq k'$  then  $\text{chop}(k) \neq \text{chop}(k')$ , and  $k$  can be derived from  $\text{chop}(k)$  and  $t$ . The chop operation also has the useful property that a single-word input key yields a single-word output sequence:  $\text{chop}(k) = \langle k \rangle$ .

With the chop function in hand, we specify the wee hash function  $h_{a,b,t,r}(k)$  for an input  $k$  of length  $t$  bits as follows:

$$h_{a,b,t,r}(k) = \text{WEE}(k, a, b, t, r, m),$$

where the procedure WEE defined on the facing page iterates through the elements of the  $w$ -bit words returned by  $\text{chop}(k)$ , applying  $f_a^r$  to the sum of the current word  $k_i$  and the previously computed hash value so far, finally returning the result obtained modulo  $m$ . This definition for variable-length and long (multiple-word) inputs is a consistent extension of the definition in equation (11.7) for short (single-word) inputs. For practical use, we recommend that  $a$  be a randomly chosen odd  $w$ -bit word,  $b$  be a randomly chosen  $w$ -bit word, and that  $r = 4$ .

Note that the wee hash function is really a hash function family, with individual hash functions determined by parameters  $a, b, t, r$ , and  $m$ . The (approximate) 5-independence of the wee hash function family for variable-length inputs can be argued based on the assumption that the 1-word wee hash function is a random oracle and on the security of the cipher-block-chaining message authentication code (CBC-MAC), as studied by Bellare et al. [42]. The case here is actually simpler than that studied in the literature, since if two messages have different lengths  $t$  and  $t'$ , then their “keys” are different:  $a + 2t \neq a + 2t'$ . We omit the details.

```

WEE( $k, a, b, t, r, m$ )
1   $u = \lceil t/w \rceil$ 
2   $\langle k_1, k_2, \dots, k_u \rangle = \text{chop}(k)$ 
3   $q = b$ 
4  for  $i = 1$  to  $u$ 
5       $q = f_{a+2t}^{(r)}(k_i + q)$ 
6  return  $q \bmod m$ 

```

This definition of a cryptographically inspired hash-function family is meant to be realistic, yet only illustrative, and many variations and improvements are possible. See the chapter notes for suggestions.

In summary, we see that when the memory system is hierarchical, it becomes advantageous to use linear probing (a special case of double hashing), since successive probes tend to stay in the same cache block. Furthermore, hash functions that can be implemented using only the computer's fast registers are exceptionally efficient, so they can be quite complex and even cryptographically inspired, providing the high degree of independence needed for linear probing to work most efficiently.

### Exercises

#### ★ 11.5-1

Complete the argument that for any odd positive integer  $a$  and any integer  $r \geq 0$ , the function  $f_a^{(r)}$  is one-to-one. Use a proof by contradiction and make use of the fact that the function  $f_a$  works modulo  $2^w$ .

#### ★ 11.5-2

Argue that a random oracle is 5-independent.

#### ★ 11.5-3

Consider what happens to the value  $f_a^{(r)}(k)$  as we flip a single bit  $k_i$  of the input value  $k$ , for various values of  $r$ . Let  $k = \sum_{i=0}^{w-1} k_i 2^i$  and  $g_a(k) = \sum_{j=0}^{w-1} b_j 2^j$  define the bit values  $k_i$  in the input (with  $k_0$  the least-significant bit) and the bit values  $b_j$  in  $g_a(k) = (2k^2 + ak) \bmod 2^w$  (where  $g_a(k)$  is the value that, when its halves are swapped, becomes  $f_a(k)$ ). Suppose that flipping a single bit  $k_i$  of the input  $k$  may cause any bit  $b_j$  of  $g_a(k)$  to flip, for  $j \geq i$ . What is the least value of  $r$  for which flipping the value of any single bit  $k_i$  may cause *any* bit of the output  $f_a^{(r)}(k)$  to flip? Explain.

---

**Problems**
**11-1 Longest-probe bound for hashing**

Suppose you are using an open-addressed hash table of size  $m$  to store  $n \leq m/2$  items.

- a.* Assuming independent uniform permutation hashing, show that for  $i = 1, 2, \dots, n$ , the probability is at most  $2^{-p}$  that the  $i$ th insertion requires strictly more than  $p$  probes.
- b.* Show that for  $i = 1, 2, \dots, n$ , the probability is  $O(1/n^2)$  that the  $i$ th insertion requires more than  $2 \lg n$  probes.

Let the random variable  $X_i$  denote the number of probes required by the  $i$ th insertion. You have shown in part (b) that  $\Pr\{X_i > 2 \lg n\} = O(1/n^2)$ . Let the random variable  $X = \max\{X_i : 1 \leq i \leq n\}$  denote the maximum number of probes required by any of the  $n$  insertions.

- c.* Show that  $\Pr\{X > 2 \lg n\} = O(1/n)$ .
- d.* Show that the expected length  $E[X]$  of the longest probe sequence is  $O(\lg n)$ .

**11-2 Searching a static set**

You are asked to implement a searchable set of  $n$  elements in which the keys are numbers. The set is static (no INSERT or DELETE operations), and the only operation required is SEARCH. You are given an arbitrary amount of time to preprocess the  $n$  elements so that SEARCH operations run quickly.

- a.* Show how to implement SEARCH in  $O(\lg n)$  worst-case time using no extra storage beyond what is needed to store the elements of the set themselves.
- b.* Consider implementing the set by open-address hashing on  $m$  slots, and assume independent uniform permutation hashing. What is the minimum amount of extra storage  $m - n$  required to make the average performance of an unsuccessful SEARCH operation be at least as good as the bound in part (a)? Your answer should be an asymptotic bound on  $m - n$  in terms of  $n$ .

**11-3 Slot-size bound for chaining**

Given a hash table with  $n$  slots, with collisions resolved by chaining, suppose that  $n$  keys are inserted into the table. Each key is equally likely to be hashed to each slot. Let  $M$  be the maximum number of keys in any slot after all the keys have



been inserted. Your mission is to prove an  $O(\lg n / \lg \lg n)$  upper bound on  $E[M]$ , the expected value of  $M$ .

- a.* Argue that the probability  $Q_k$  that exactly  $k$  keys hash to a particular slot is given by

$$Q_k = \left(\frac{1}{n}\right)^k \left(1 - \frac{1}{n}\right)^{n-k} \binom{n}{k}.$$

- b.* Let  $P_k$  be the probability that  $M = k$ , that is, the probability that the slot containing the most keys contains  $k$  keys. Show that  $P_k \leq n Q_k$ .
- c.* Show that  $Q_k < e^k / k^k$ . *Hint:* Use Stirling's approximation, equation (3.25) on page 67.
- d.* Show that there exists a constant  $c > 1$  such that  $Q_{k_0} < 1/n^3$  for  $k_0 = c \lg n / \lg \lg n$ . Conclude that  $P_k < 1/n^2$  for  $k \geq k_0 = c \lg n / \lg \lg n$ .
- e.* Argue that

$$E[M] \leq \Pr \left\{ M > \frac{c \lg n}{\lg \lg n} \right\} \cdot n + \Pr \left\{ M \leq \frac{c \lg n}{\lg \lg n} \right\} \cdot \frac{c \lg n}{\lg \lg n}.$$

Conclude that  $E[M] = O(\lg n / \lg \lg n)$ .

#### 11-4 Hashing and authentication

Let  $\mathcal{H}$  be a family of hash functions in which each hash function  $h \in \mathcal{H}$  maps the universe  $U$  of keys to  $\{0, 1, \dots, m-1\}$ .

- a.* Show that if the family  $\mathcal{H}$  of hash functions is 2-independent, then it is universal.
- b.* Suppose that the universe  $U$  is the set of  $n$ -tuples of values drawn from  $\mathbb{Z}_p = \{0, 1, \dots, p-1\}$ , where  $p$  is prime. Consider an element  $x = \langle x_0, x_1, \dots, x_{n-1} \rangle \in U$ . For any  $n$ -tuple  $a = \langle a_0, a_1, \dots, a_{n-1} \rangle \in U$ , define the hash function  $h_a$  by

$$h_a(x) = \left( \sum_{j=0}^{n-1} a_j x_j \right) \bmod p.$$

Let  $\mathcal{H} = \{h_a : a \in U\}$ . Show that  $\mathcal{H}$  is universal, but not 2-independent. (*Hint:* Find a key for which all hash functions in  $\mathcal{H}$  produce the same value.)

- c. Suppose that we modify  $\mathcal{H}$  slightly from part (b): for any  $a \in U$  and for any  $b \in \mathbb{Z}_p$ , define

$$h'_{ab}(x) = \left( \sum_{j=0}^{n-1} a_j x_j + b \right) \bmod p$$

and  $\mathcal{H}' = \{h'_{ab} : a \in U \text{ and } b \in \mathbb{Z}_p\}$ . Argue that  $\mathcal{H}'$  is 2-independent. (*Hint*: Consider fixed  $n$ -tuples  $x \in U$  and  $y \in U$ , with  $x_i \neq y_i$  for some  $i$ . What happens to  $h'_{ab}(x)$  and  $h'_{ab}(y)$  as  $a_i$  and  $b$  range over  $\mathbb{Z}_p$ ?)

- d. Alice and Bob secretly agree on a hash function  $h$  from a 2-independent family  $\mathcal{H}$  of hash functions. Each  $h \in \mathcal{H}$  maps from a universe of keys  $U$  to  $\mathbb{Z}_p$ , where  $p$  is prime. Later, Alice sends a message  $m$  to Bob over the internet, where  $m \in U$ . She authenticates this message to Bob by also sending an authentication tag  $t = h(m)$ , and Bob checks that the pair  $(m, t)$  he receives indeed satisfies  $t = h(m)$ . Suppose that an adversary intercepts  $(m, t)$  en route and tries to fool Bob by replacing the pair  $(m, t)$  with a different pair  $(m', t')$ . Argue that the probability that the adversary succeeds in fooling Bob into accepting  $(m', t')$  is at most  $1/p$ , no matter how much computing power the adversary has, even if the adversary knows the family  $\mathcal{H}$  of hash functions used.

---

## Chapter notes

The books by Knuth [261] and Gonnet and Baeza-Yates [193] are excellent references for the analysis of hashing algorithms. Knuth credits H. P. Luhn (1953) for inventing hash tables, along with the chaining method for resolving collisions. At about the same time, G. M. Amdahl originated the idea of open addressing. The notion of a random oracle was introduced by Bellare et al. [43]. Carter and Wegman [80] introduced the notion of universal families of hash functions in 1979.

Dietzfelbinger et al. [113] invented the multiply-shift hash function and gave a proof of Theorem 11.5. Thorup [437] provides extensions and additional analysis. Thorup [438] gives a simple proof that linear probing with 5-independent hashing takes constant expected time per operation. Thorup also describes the method for deletion in a hash table using linear probing.

Fredman, Komlós, and Szemerédi [154] developed a perfect hashing scheme for static sets—“perfect” because all collisions are avoided. An extension of their method to dynamic sets, handling insertions and deletions in amortized expected time  $O(1)$ , has been given by Dietzfelbinger et al. [114].

The wee hash function is based on the RC6 encryption algorithm [379]. Leiser et al. [292] propose an “RC6MIX” function that is essentially the same as the

wee hash function. They give experimental evidence that it has good randomness, and they also give a “DOTMIX” function for dealing with variable-length inputs. Bellare et al. [42] provide an analysis of the security of the cipher-block-chaining message authentication code. This analysis implies that the wee hash function has the desired pseudorandomness properties.