RDBMS-3

Indexing + ACID

DCC → ACID

RAM → volatile / non persist

HDD/SSD → non volatile / persistent

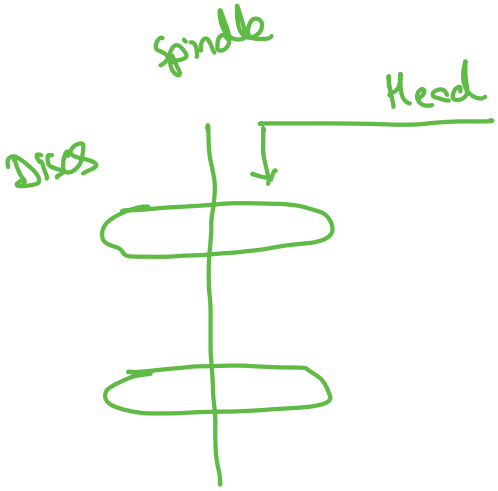

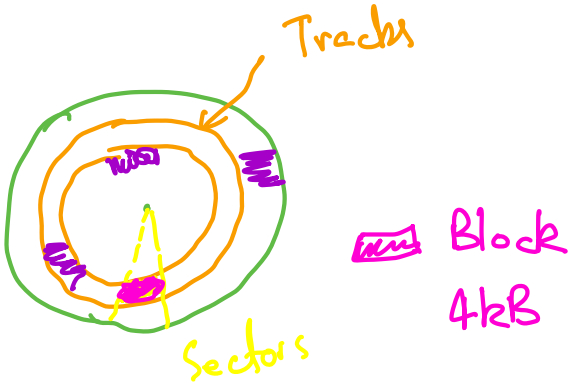

db

sid sAddress sCity sSubject 1kB

⋮

10,000 rows

db size = 10,000 kB ~ 10 MB

Each block ~ 4kB

Total = 10,000 kB

⇒ No of Blocks = 2500 blocks → Spread Across

Q → Give me all ddc for student id = 120.

sid ⟶ Block X → 120 → Block X

Indexing

I Simplest

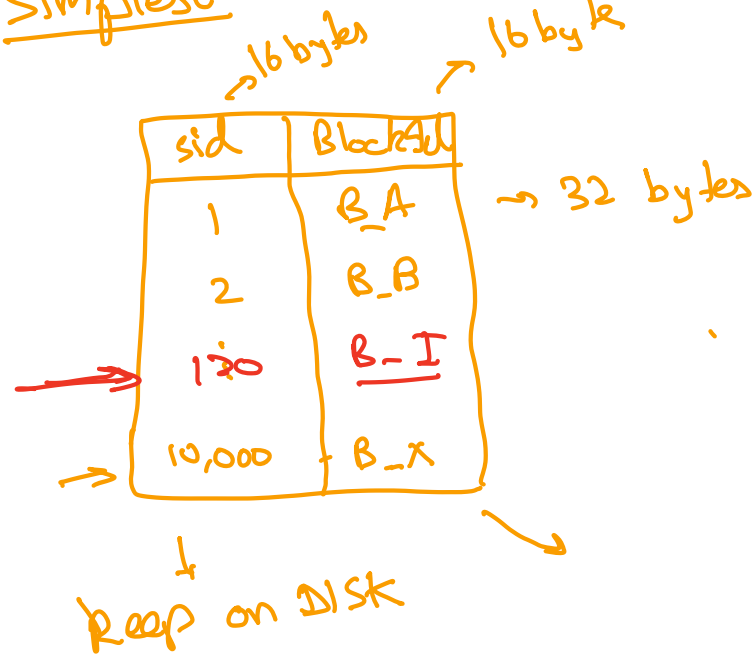

- Each block → 4kB
  Each row = 32
  Each Block = $\frac{4000}{32}$ rows
  ⟶ = ~128

RAM size = 8 kB

Index table = $\frac{10000}{128}$

= 78 blocks

<sid = 120>

Total Disk blocks reads = 78 + 1 → to get data from block

↓

to find correct row in Index

L = 0

R = 10,000

→ M = $\frac{L+R}{2}$

Find Index 120

↓

120 = M ?

L = 0

R = 10,000

→ M = 500

Is 120 sid is in whichever block 500 is in?

Multi level Index

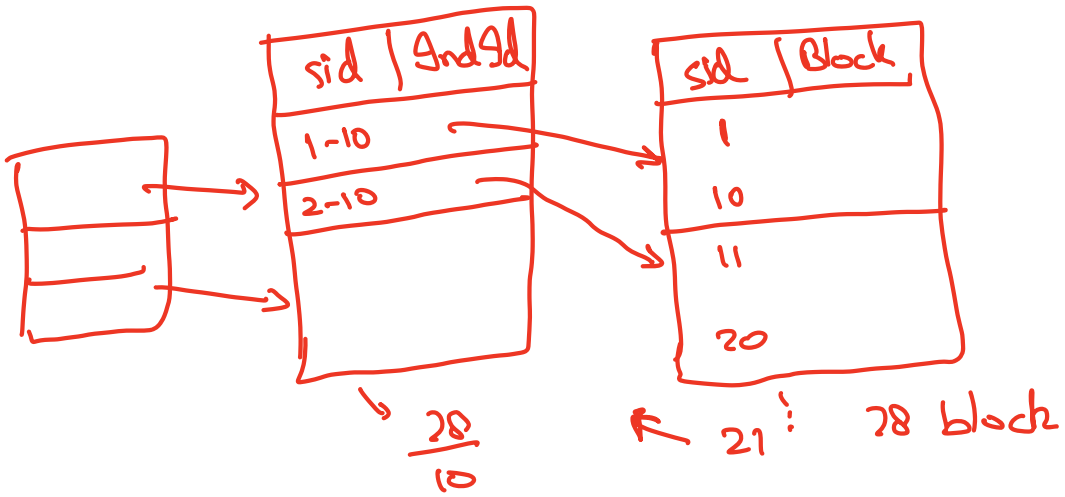

$\frac{78}{10}$ → 8 blocks ← 21 : 78 block

$\log_b(78)$

Blocks overhead

→ Update or add index

II BST

(1 row = 1 block)

↳ node in tree

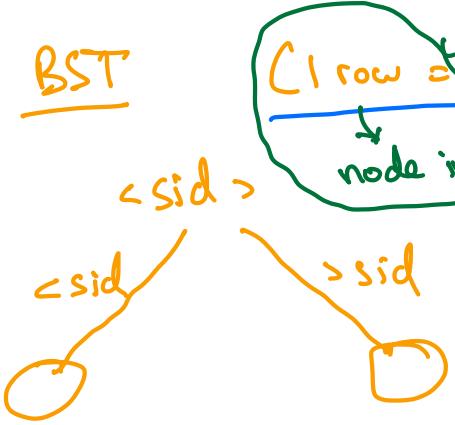

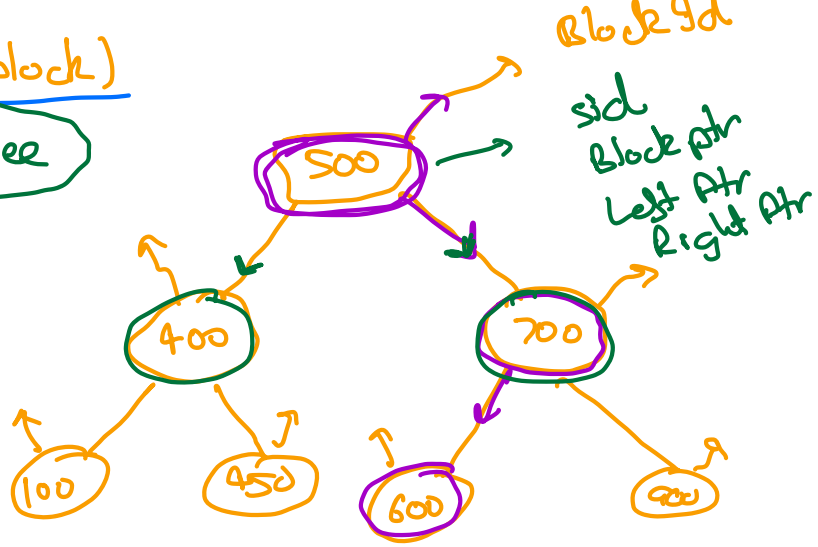

Total n nodes in tree

↳ log (n) → number of rows = number of blocks

Problems

1. In each node, storing very less. ⇒ Way more blocks

III M-way Tree → Showed is possible!

→ BST with multiples key & multiple children

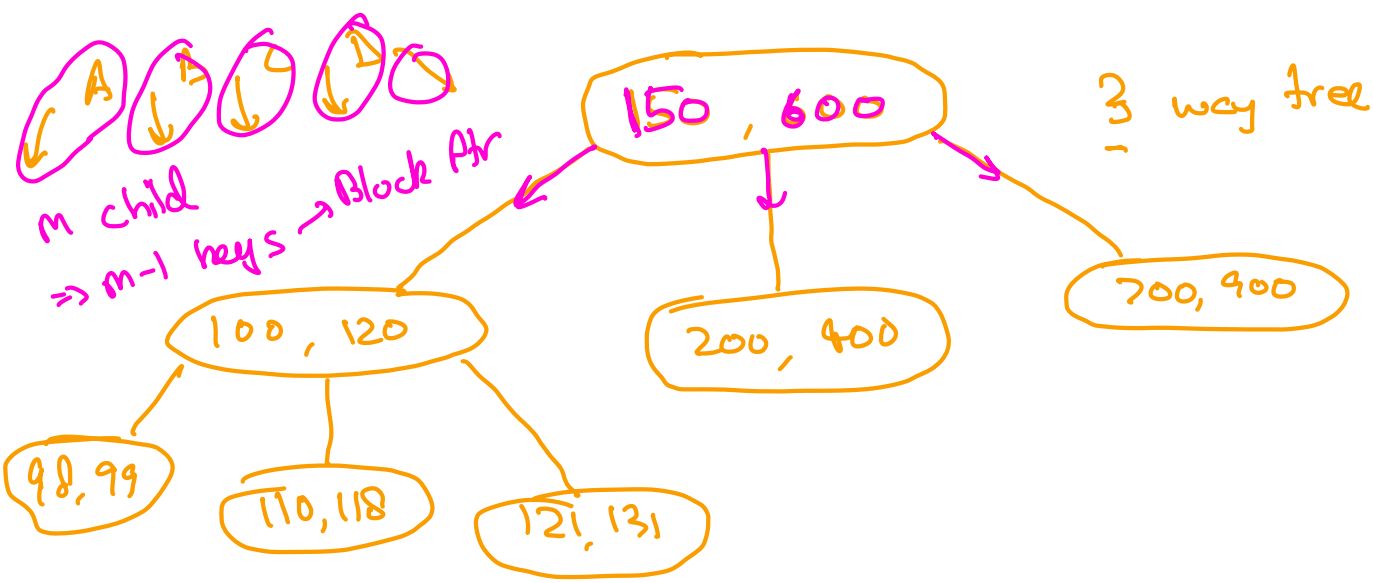

Each node will have → m-1 keys ⇒ (m-1) Block ptr

m child pointers

Size of each node ≤ Block size

10,000 rows → 1 row per node = 10k nodes

→ 10 row ... = 1k nodes

5 rows = 200 nodes

Size of each node = (m-1) x Size of key

+ (m-1) x Size of Block Pointer

+ m x Size of Child Ptr

≤ Size of Block

$(m-1) \times 16 + (m-1) \times 16 + m \times 16 \leq 4kB$

↓

$m \leq 84$

B Tree

→ self balancing M way tree

1. Each node must contain at most m children
2. Every non leaf & non root node must have at least m/2 children
3. Root must have at least 2 children
4. Trees are created in bottom up

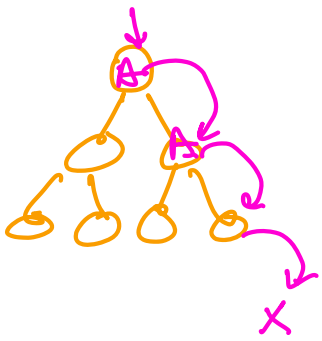

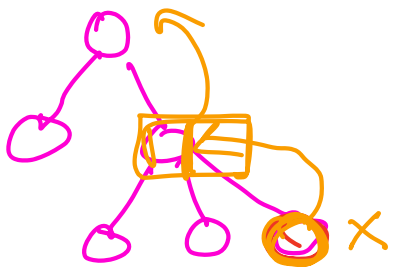

Insertion

10, 20, 30, 40, 50, 60, 70, 80

Degree ⇒ 3 keys

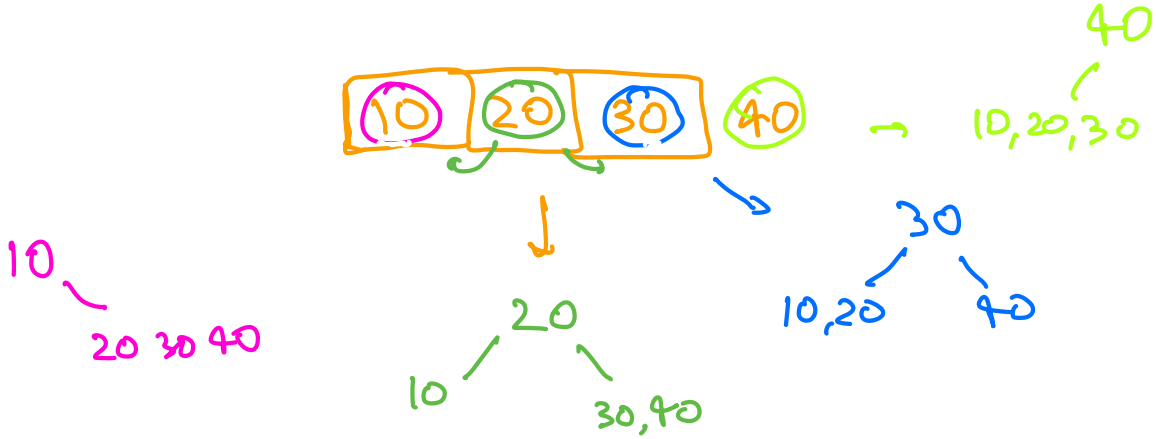

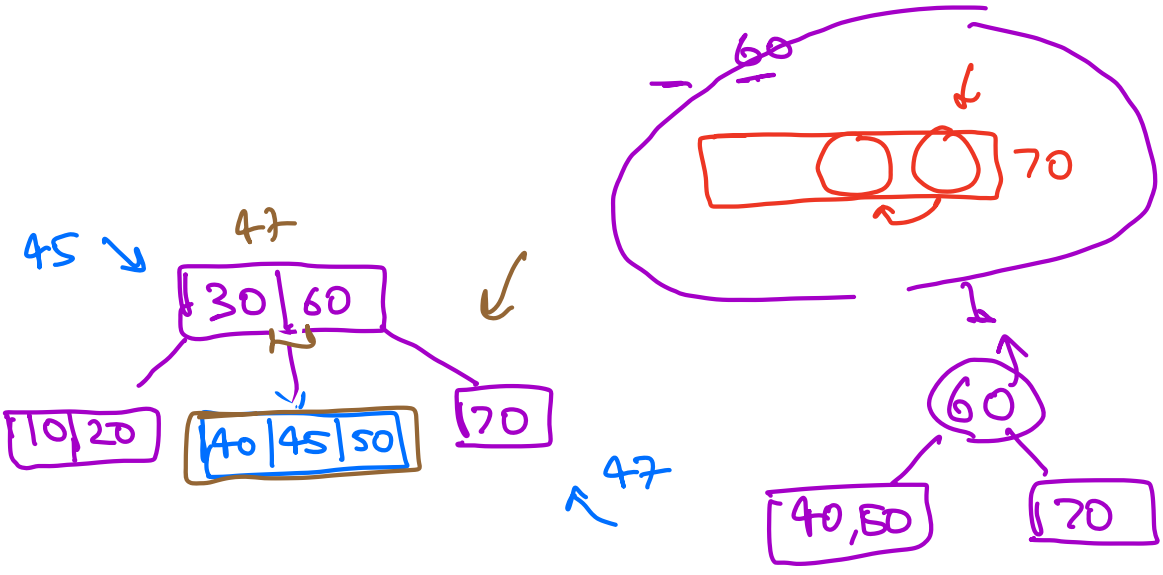

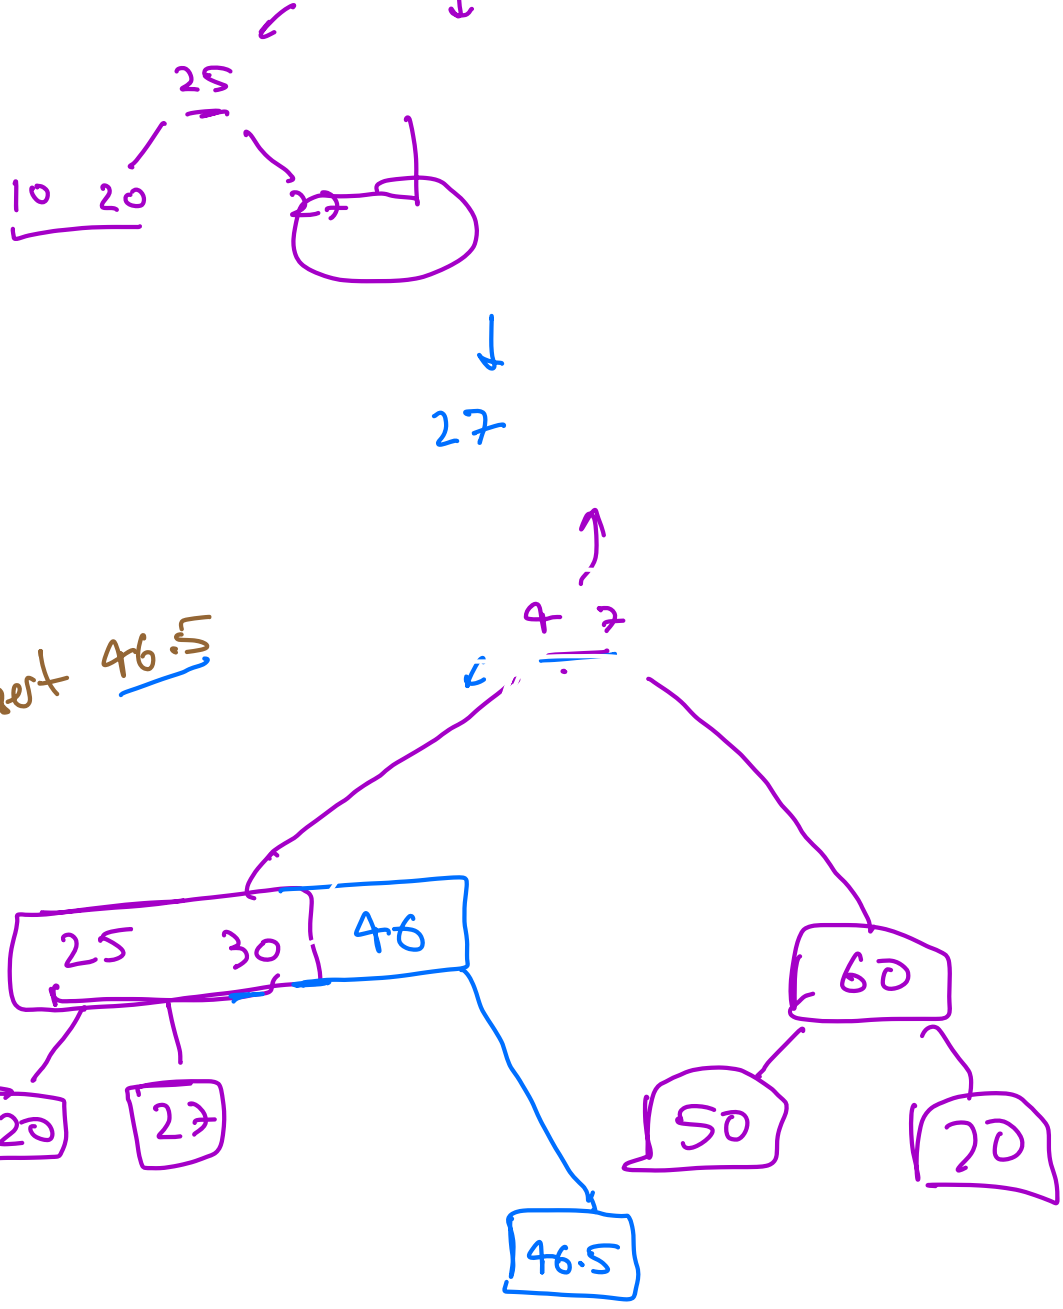
25
10 20
27
27
Insert 46.5
47
25 30 46
60
10 20
27
50
70
46.5

Problem

Range Queries → Give me all indexes → 35 – 50

Blocks reads = $(R - L) \times \log(n)$

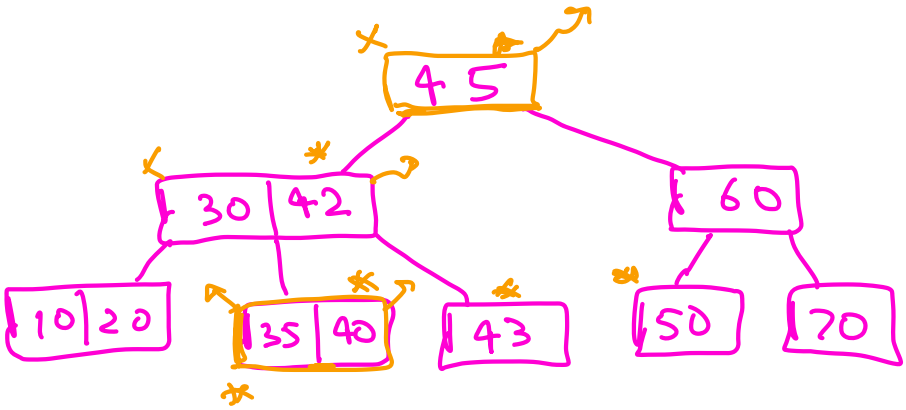

B + Trees

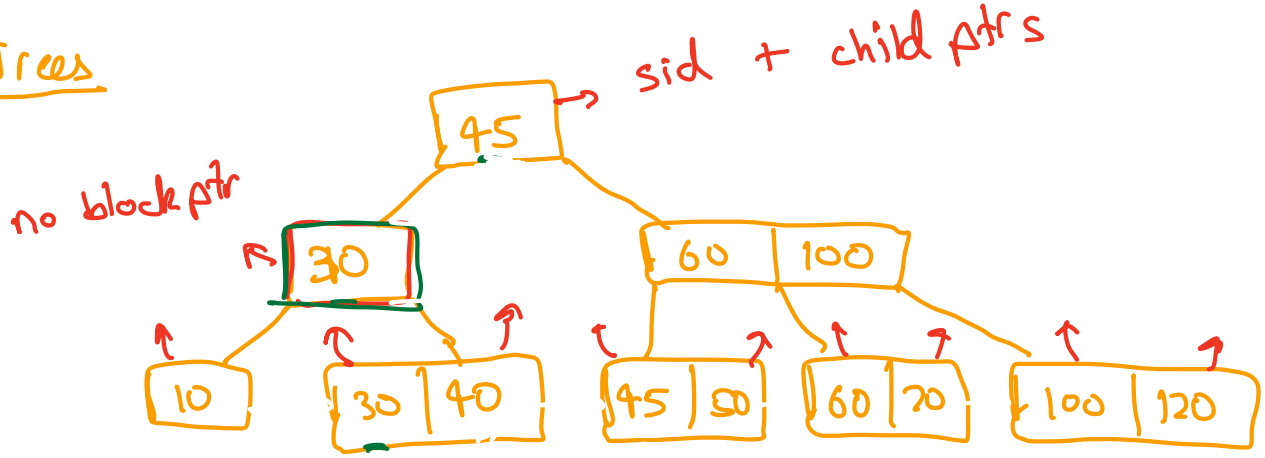

→ either entire T completes properly or it has no effect

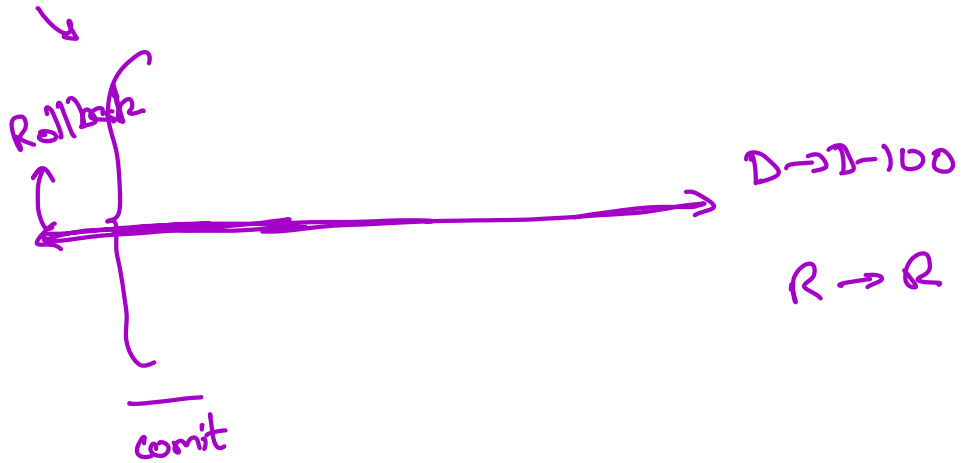

Consistency

Balance ≥ 0 → Data Integrity

either execution will follow consistency or fail

D → 200

E → D = D-100 ✓

D = D-300

fail

Isolation → guarantees that multiples atomic executions running at the same time will not impact each other

D = 500 → 400

R = 1000 → 1100 → 900

A = 1000 → 1100

D ⇒ R , R ⇒ A(c)y

↓

✓ get (D) → D' = 500

✓ get (R) → R' = 1000

✓ updde D = D'-100

✓ updde R = R'+100 → 1100

✓ get (R) → R'' = 1000

✓ get (A) → A' = 1000

✓ updde R = R''-100 →

updde A = A'+100 →

Durability

↳ Power loss / System Crash / malfunction etc → data will persist